

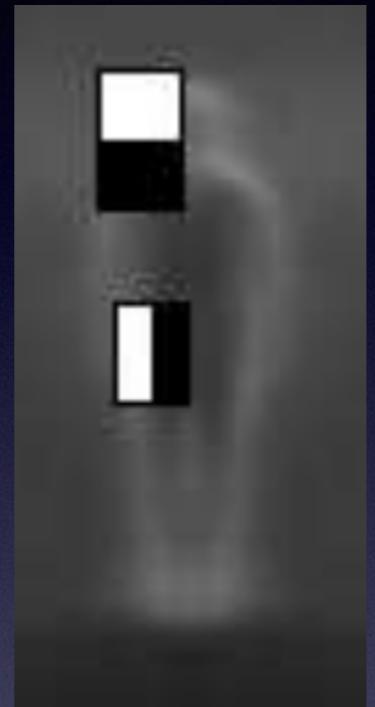
# Deformable Parts Model

Carlo Tomasi

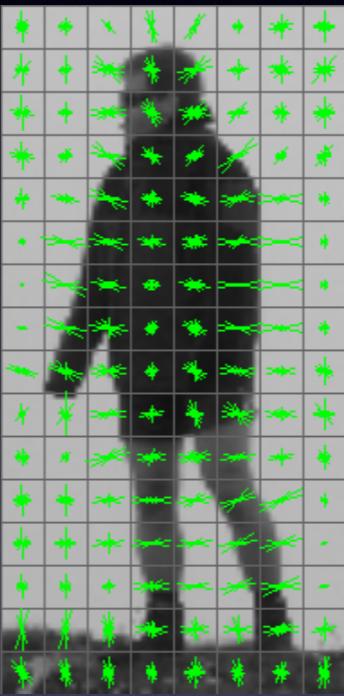
# Models for Person Detection



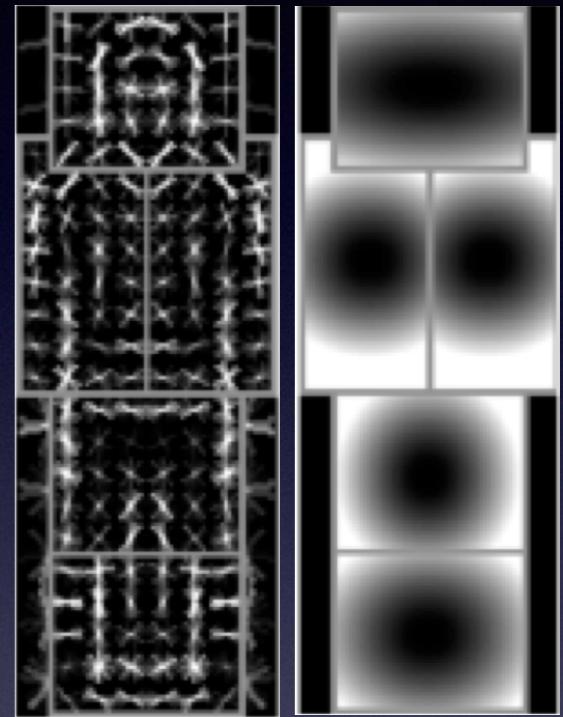
picture by Li, Fergus, Torralba



picture by Mekonnen, Lerasle, Herboulot



picture by Jürgen Brauer



picture by Felzenszwalb, Girshick, McAllester, Ramanan

bag of features:  
no shape model

Sivic *et al.* 2003  
Csurka *et al.* 2004

...

sparse features:  
fixed constellation  
of Haar features  
Viola & Jones 2001

...

grid of features:  
histograms of  
gradients  
Dalal & Triggs 2005

...

deformable parts:  
flexible constellation  
of HOG features  
Felzenszwalb *et al.* 2008

...

# Trade-Offs

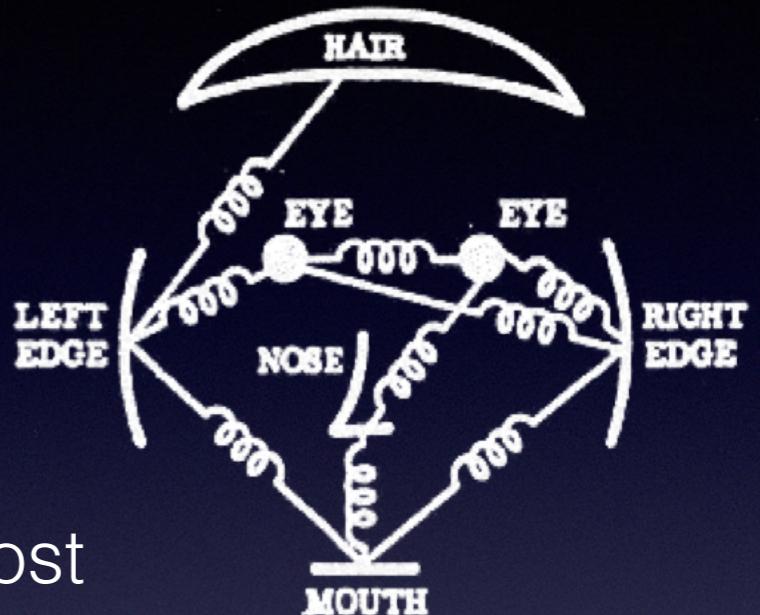
- All are sliding-window methods: expensive but embarrassingly parallel
- Bags of features: general, simple, but no shape
- Sparse features and feature grids: simple, but “people as popsicles”
- Deformable parts: accounts for body articulation, but more expensive to train and run

# Deformable Parts Model

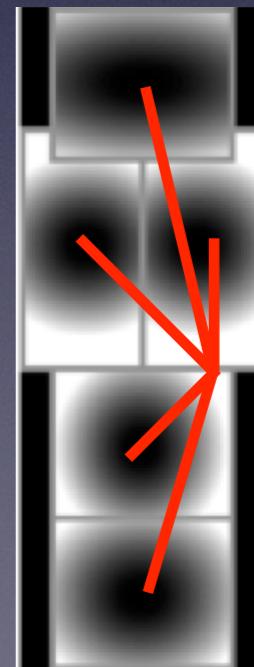
- An old idea: Fischler and Elschlager, *Pictorial Structures*, 1973

$$\{\hat{\mathbf{x}}_p\} = \arg \max_{\{\mathbf{x}_p\}} \sum_p f(\mathbf{x}_p) - \sum_{p,q} d(\mathbf{x}_p, \mathbf{x}_q)$$

LEFT  
EDGE

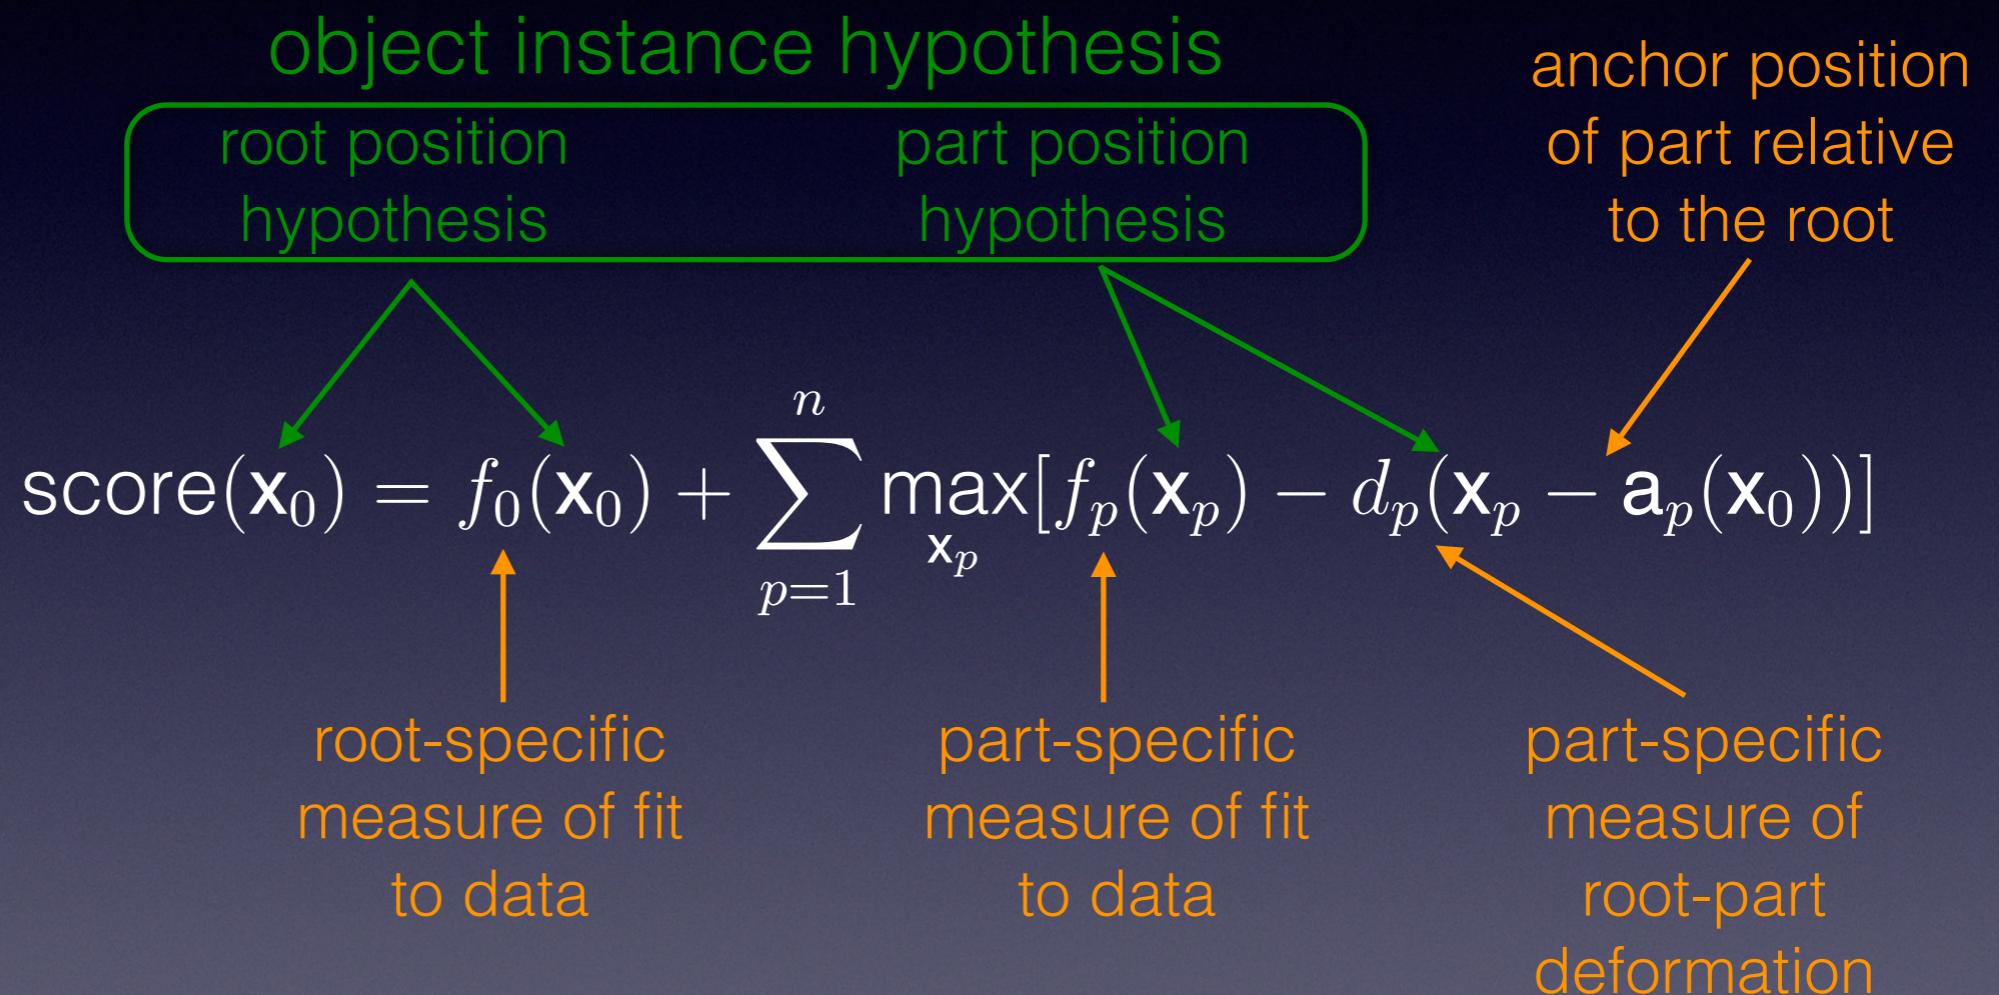
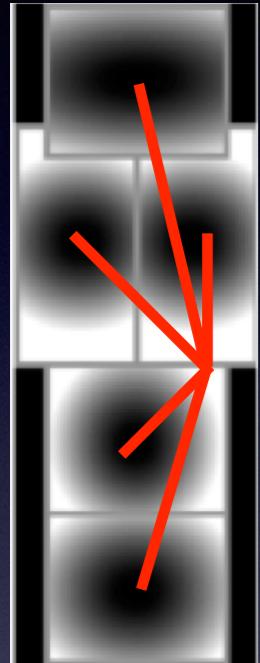


- Key difficulty is combinatorially explosive matching complexity during detection
  - Solution [Felzenszwalb & Huttenlocher 2000; Felzenszwalb, Girshick, McAllester, Ramanan, 2010]: any pair  $\rightarrow$  star graph
  - Use dynamic programming



# Deformable Parts Model

root: the whole window; parts: subwindows



- object detections are strong local maxima of  $\text{score}(\mathbf{x}_0)$
- corresponding  $\mathbf{x}_p$  yield model-to-instance correspondence
- mixtures of DPMs handle large intra-class variations

# Fit and Deformation Measures

image features



$$f_p(\mathbf{x}) = \boldsymbol{\beta}_p^T \boldsymbol{\varphi}(\mathbf{x})$$

↑  
linear function  
of features defined  
on an image pyramid

$$\boldsymbol{\eta} = [x - a_p, \ y - b_p, \ (x - a_p)^2, \ (y - b_p)^2]^T$$



$$d_p(\mathbf{x} - \mathbf{a}_p(\mathbf{x}_0)) = \boldsymbol{\delta}_p^T \boldsymbol{\eta}$$

↑  
quadratic function  
of part-anchor  
displacement

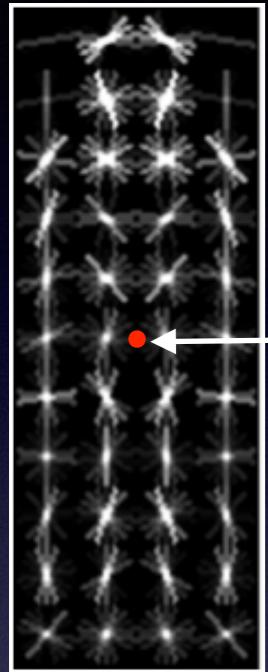
**parabola!**

Training determines model parameters

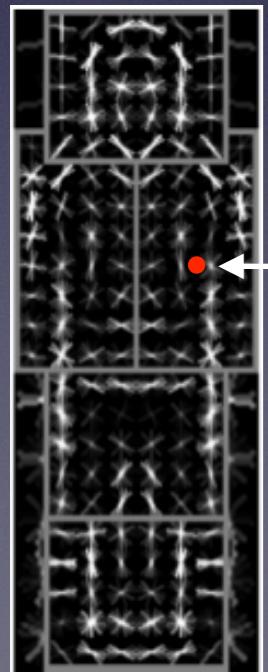
$$\mathbf{w}^T = (\boldsymbol{\beta}_0; \boldsymbol{\beta}_1, \boldsymbol{\delta}_1, \mathbf{a}_1, \dots, \boldsymbol{\beta}_n, \boldsymbol{\delta}_n, \mathbf{a}_n)$$

$\mathbf{x}_0 = (x_0, y_0, \ell_0)$  and  $\mathbf{x}_p = (x_p, y_p, \ell_0 - \lambda)$  for  $p \geq 1$  are instance parameters

# Features



$$\mathbf{x}_0 = (x_0, y_0, \ell_0)$$



twice the resolution

$$\mathbf{x}_p = (x_p, y_p, \ell_0 - \lambda) \text{ for } p \geq 1$$

- HOG features for both root and parts  
[Dalal & Triggs 2006]
- Part HOGs are one octave finer than root HOG
- Some dimensionality reduction through PCA saves both training and detection complexity
- All features computed on a fine image pyramid for scale sensitivity

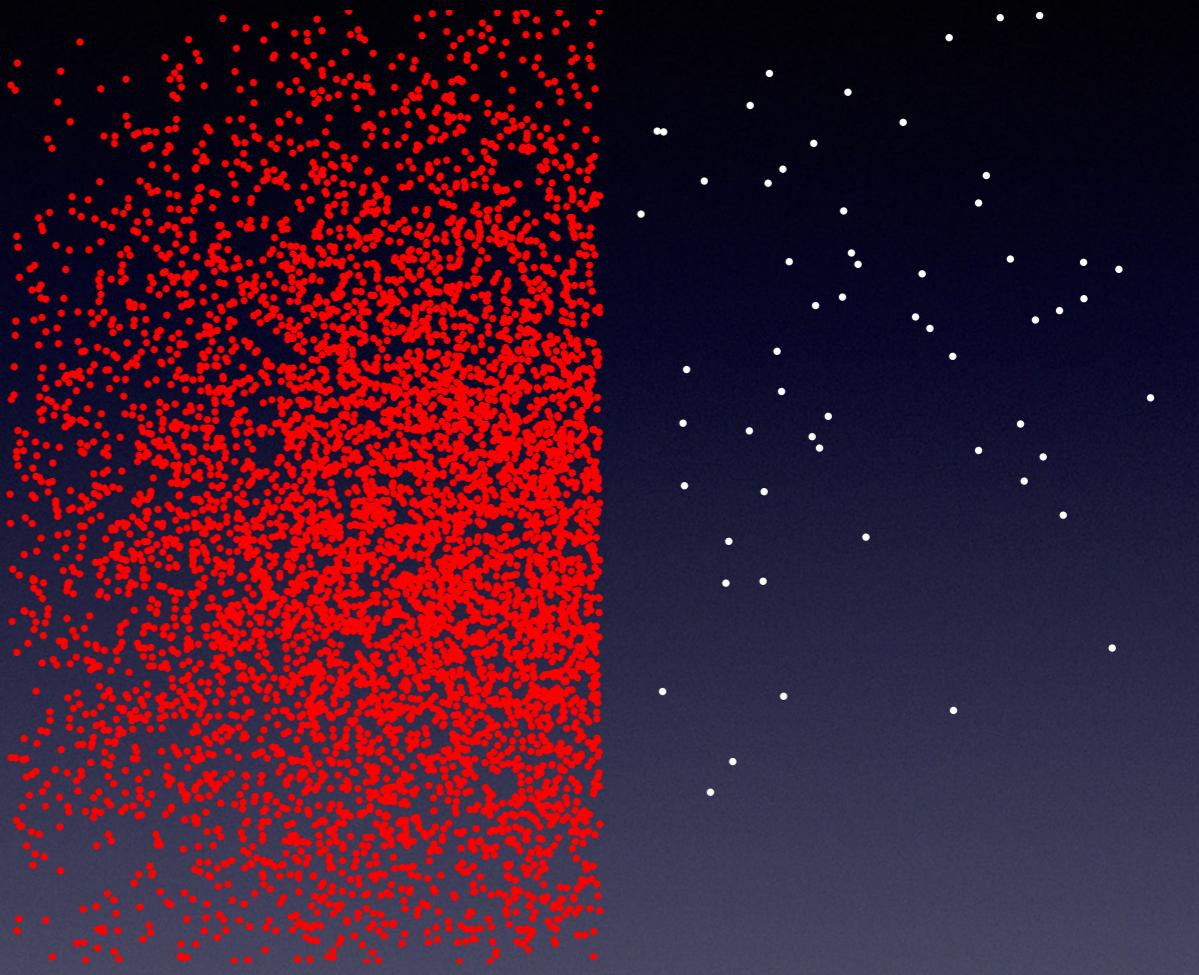
# Training

- Standard SVM classifier:  $y = \text{sign}(\mathbf{w}^T \mathbf{f} - b)$
- Latent SVM classifier:  $y = \text{sign} \max_{\mathbf{x}} [\mathbf{w}^T \mathbf{f}(\mathbf{x}) - b]$
- This new problem leads to nearly the same optimization problem as the standard soft-margin SVM:
$$\arg \min_{\mathbf{w}} \left[ \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \max\{0, 1 - y_i \max_{\mathbf{x}} [\mathbf{w}^T \mathbf{f}(\mathbf{x}) - b]\} \right]$$
- However, the optimization target is no longer convex
- Use stochastic gradient descent to optimize, but lose global convergence guarantees
- Requires careful initialization

part  
location  
hypotheses

# Choosing Negative Examples

- Many more negative than positive examples (“not a person”)
  - Using all examples would lead to slow training and many support vectors (slow detection)
  - Random subset would lead to poor representation along boundary
  - Use data mining techniques to choose “hard” negative examples



- start with random subset
- learn classifier
- collect misclassified negatives
- repeat

# Detection

$$\text{score}(\mathbf{x}_0) = f_0(\mathbf{x}_0) + \sum_{p=1}^n \max_{\mathbf{x}_p} [f_p(\mathbf{x}_p) - d_p(\mathbf{x}_p - \mathbf{a}_p(\mathbf{x}_0))]$$

$\forall \mathbf{x}$  in the pyramid, compute  $\varphi(\mathbf{x})$

$\forall p \in \{0, \dots, n\}$ ,  $\forall \mathbf{x}$  in the pyramid, compute  $f_p(\mathbf{x}) = \boldsymbol{\beta}_p^T \varphi(\mathbf{x})$

$\forall p \in \{1, \dots, n\}$ ,  $\forall \mathbf{a}$  in the pyramid, compute  $D_p(\mathbf{a}) = \max_{\mathbf{x}} [f_p(\mathbf{x}) - d_p(\mathbf{x} - \mathbf{a})]$

$\forall \mathbf{x}_0$  in the pyramid, compute  $\text{score}(\mathbf{x}_0) = f_0(\mathbf{x}_0) + \sum_{p=1}^n D_p(\mathbf{a}_p(\mathbf{x}_0))$

select high-scoring root positions by non-maximum suppression

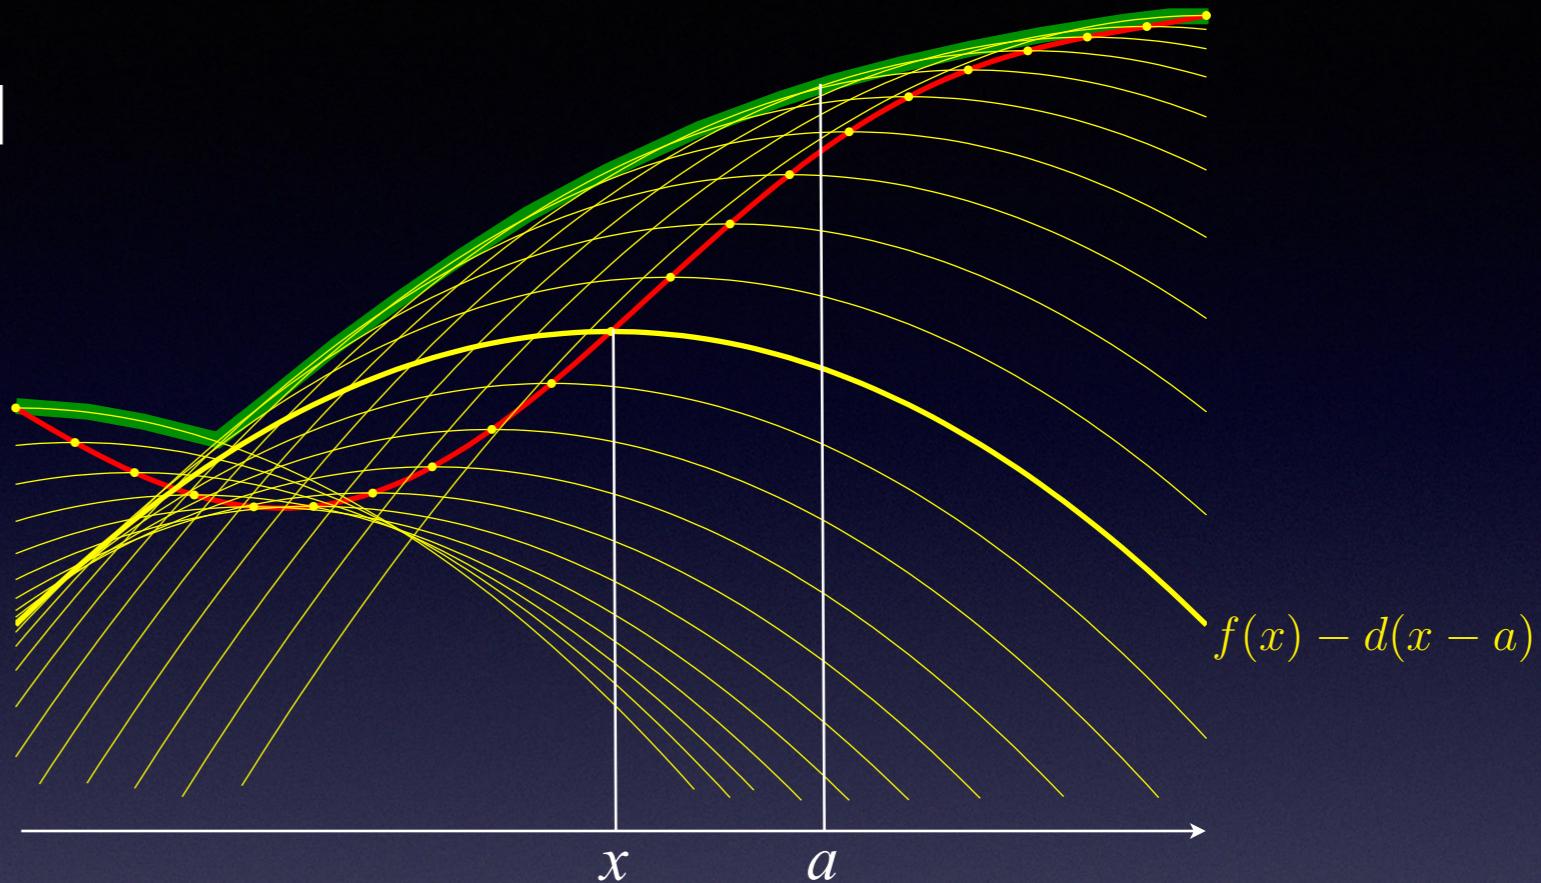
# Generalized Distance Transform

[Rosenfeld and Pfaltz 1966,  
Felzenszwalb and Huttenlocher 2004]

## One dimension:

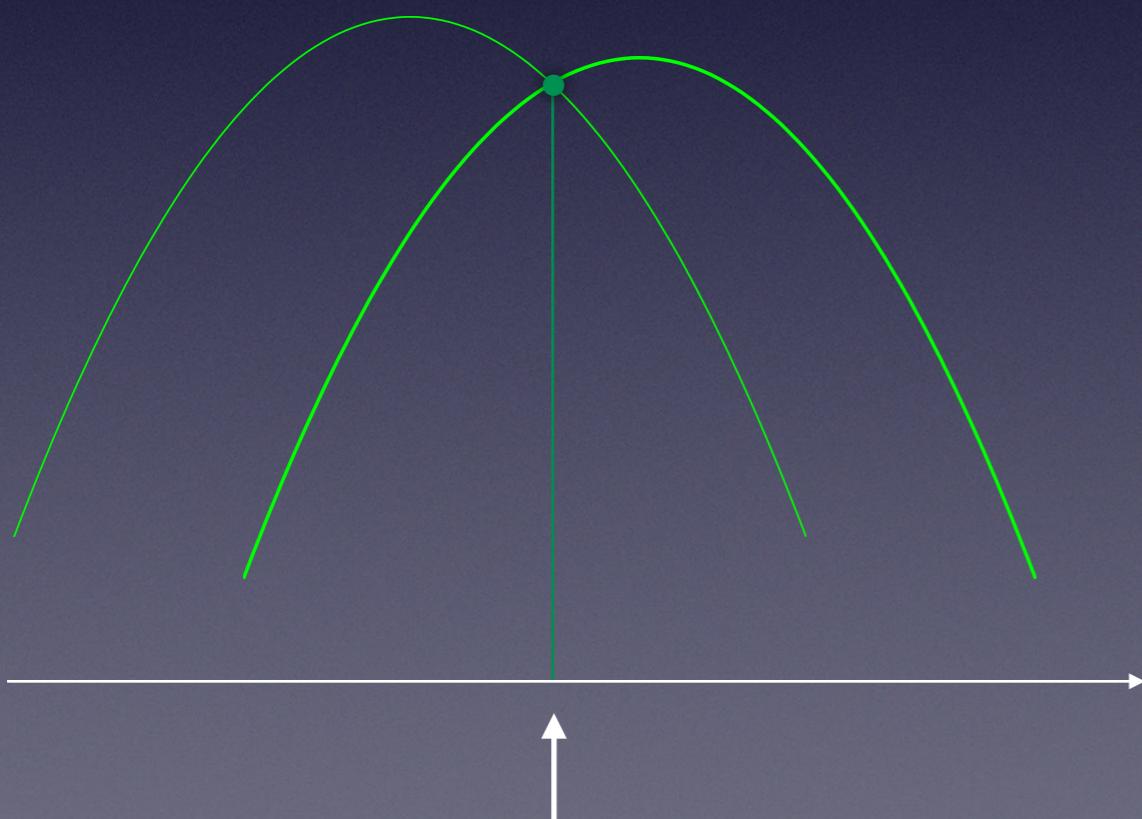
$$D(a) = \max_x [f(x) - d(x - a)]$$

A single left-to-right sweep does the job:



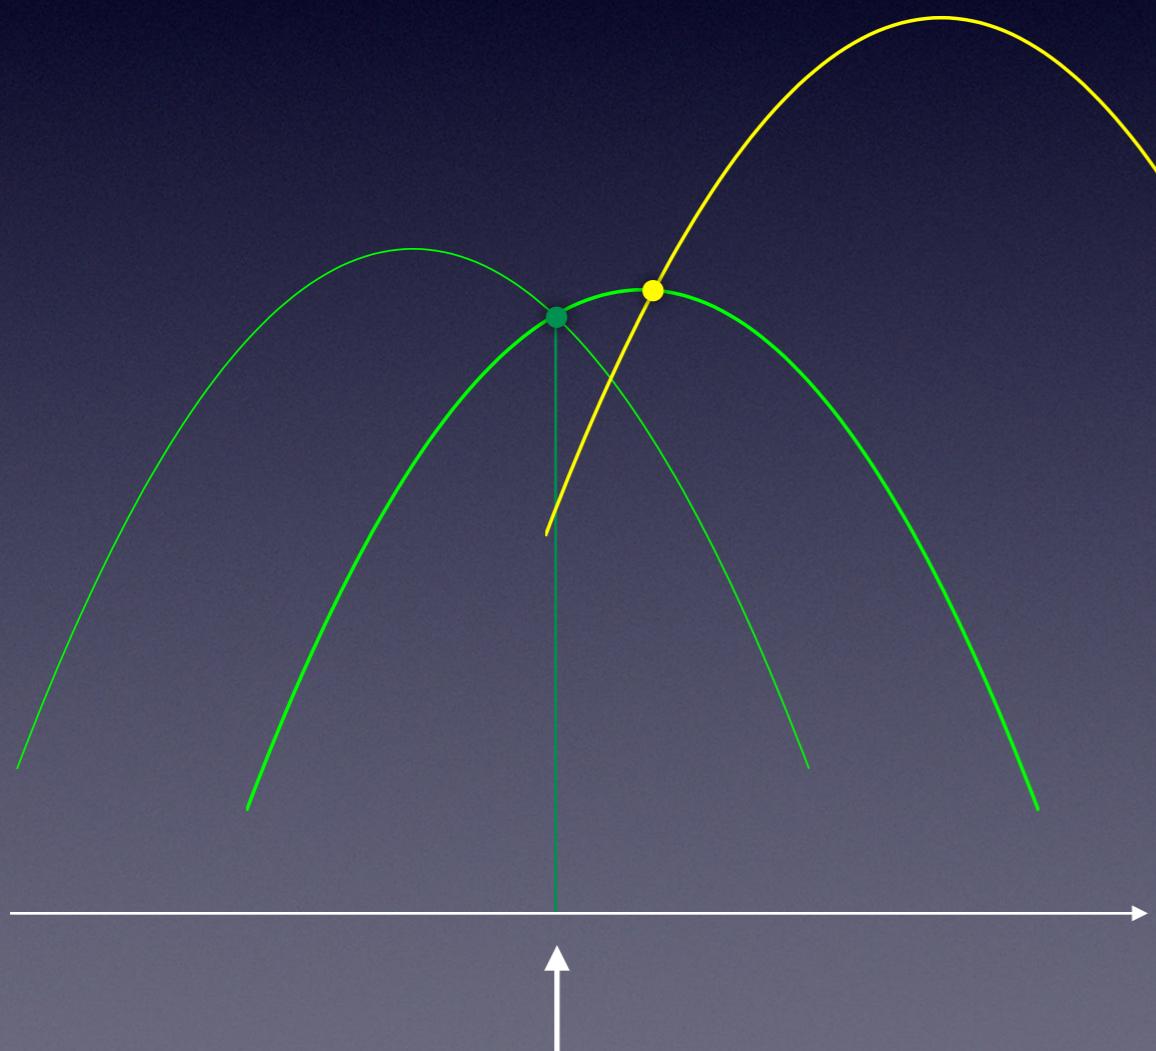
- Not all parabolas touch the upper envelope
  - Examine parabolas one at a time, left to right
  - Keep track of whether and where the new parabola intersects the last parabola in the envelope (intersections are solutions of second-degree equations)
  - Connect relevant parabola pieces to form upper envelope

# Two Cases



last parabola intersection in the current envelope

# Two Cases



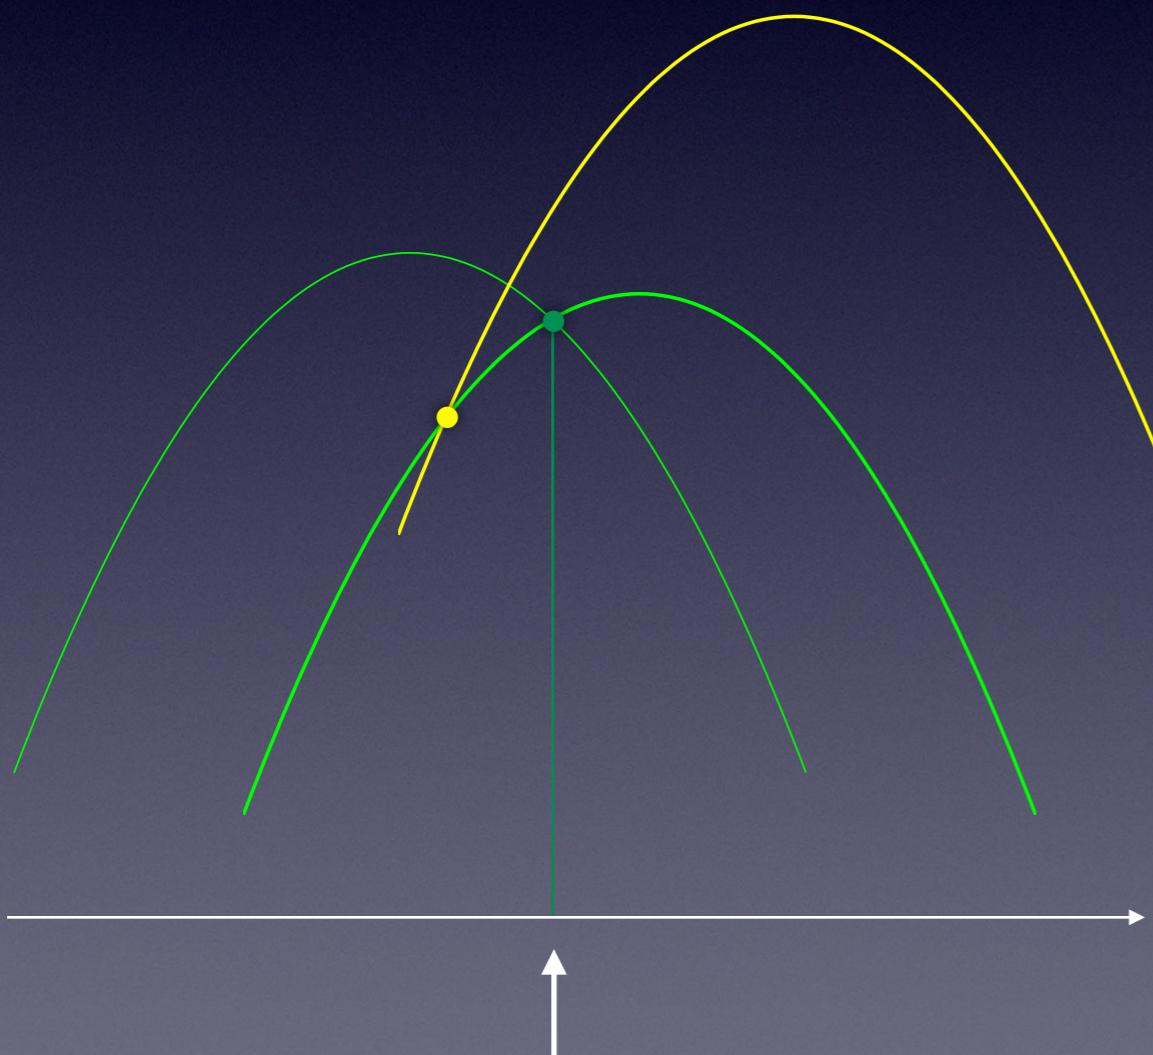
Intersection of new parabola  
and last envelope parabola  
is to the right of the  
last envelope intersection



add the new parabola  
to the envelope

last parabola intersection in the current envelope

# Two Cases



last parabola intersection in the current envelope

Intersection of new parabola  
and last envelope parabola  
is to the left of the  
last envelope intersection

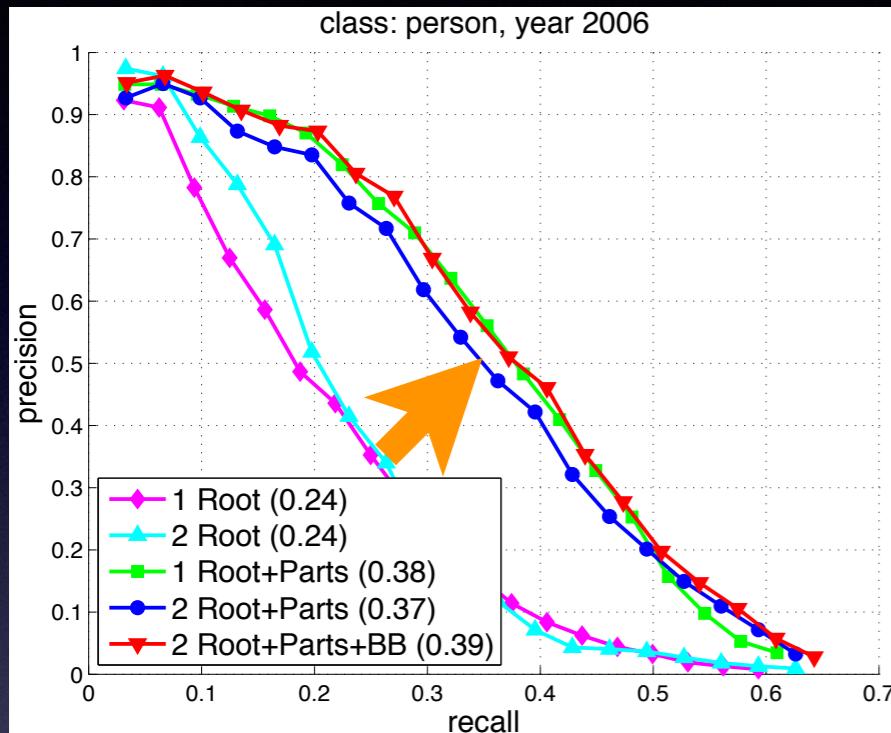


remove the last parabola  
from the envelope

# Distance Transform in 2D

$$\begin{aligned} & \max_{\mathbf{x}} [f(\mathbf{x}) - d_2(\mathbf{x} - \mathbf{a})] \\ &= \max_{(x,y)} [f(x, y) - d(x - a) - d(y - b)] \\ &= \max_x \{ \overbrace{\max_y [f(x, y) - d(y - b)]}^{\text{1D case for each } x} - d(x - a) \} \\ &= \underbrace{\max_x [D(x, b) - d(x - a)]}_{\text{1D case for each } b} \end{aligned}$$

# Performance of DPM



Effect of including parts

Average Precision(%)  
in PASCAL VOC

[DPM implementation, data, and participants change every year]

Year	DPM	Best
2005		12
2006		16
2007		22
2008	27	42
2009	36	43
2010	45	47
2011	46	52
2012	46	46

## RUNNING TIMES

4 hours to train on typical PASCAL VOC database.  
2 seconds per image on standard laptop

Zhang *et al.*,  
Inst. of Automation,  
Chinese Acad.  
Science.  
Based on DPM,  
richer part  
location model,  
some context  
information