

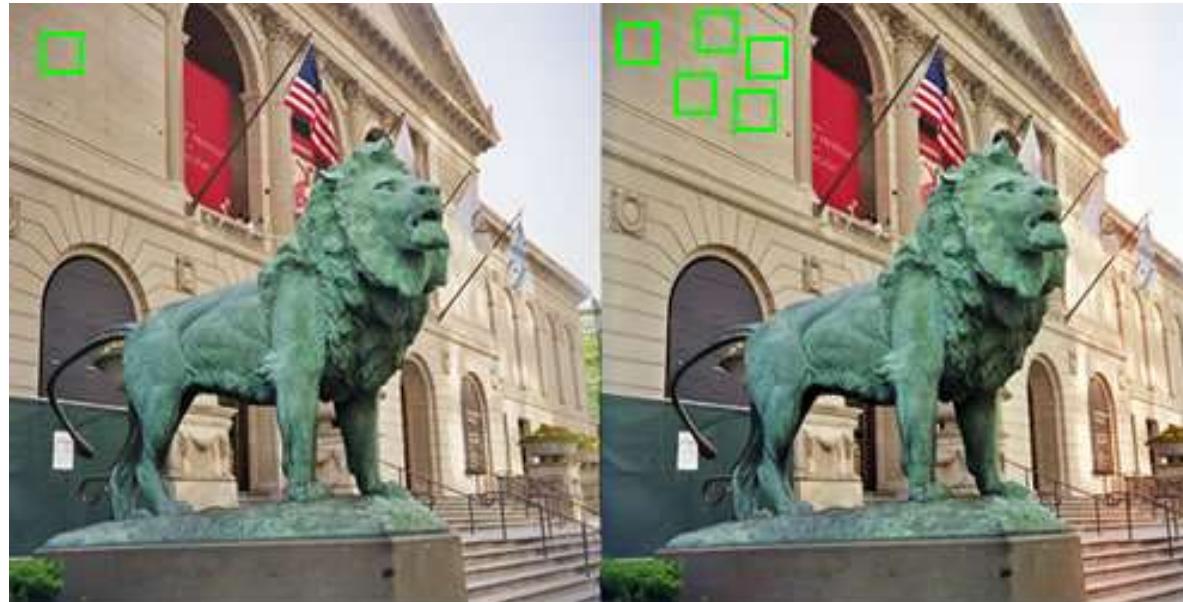
Lecture 03- Feature Points and Lines

EE382-Visual localization & Perception

Danping Zou @Shanghai Jiao Tong
University

Feature points: What are they?

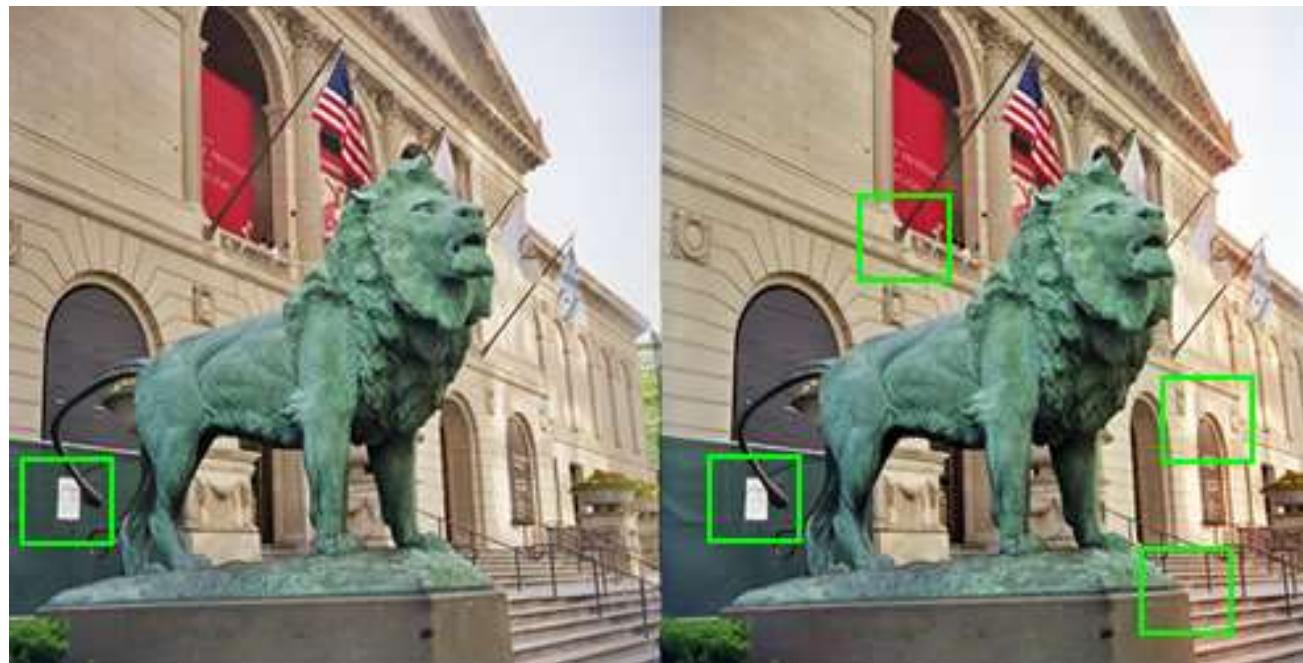
- Feature points are the backbone of a lot of computer vision algorithms. But what is a feature point?



Feature points do not locate at position with unrecognizable appearance.

Features points: What are they?

- Intuitively, a patch with *unique appearance*



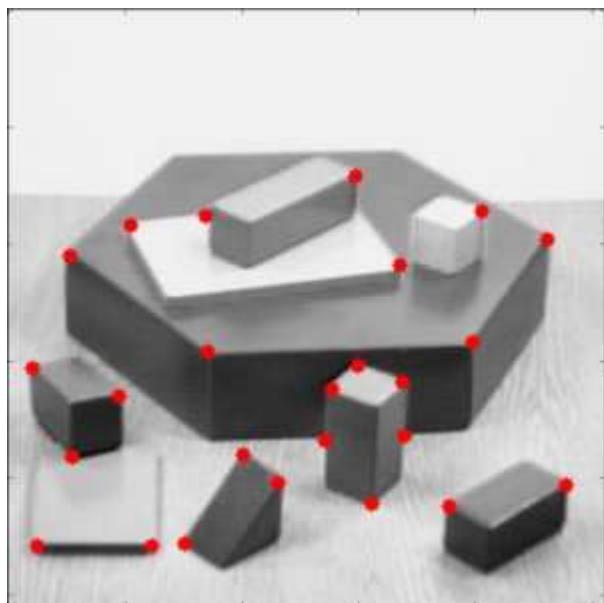
But there could be a few of such patches in an image.

Features points: What are they?

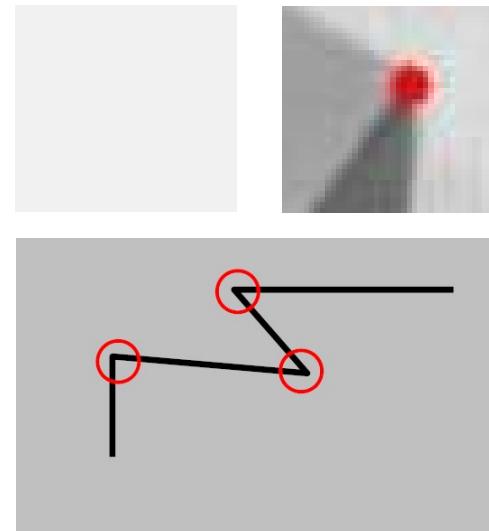
- Three properties a good feature point should have:
 - 1) Stable : The feature point can be detected repeatedly in different views.
 - 2) Recognizable : It has rich local information that make it easy to recognized.
 - 3) Accurate : Feature point can be located in sub-pixel accuracy.

Feature detection

- We introduce a classic feature point - Harris corner



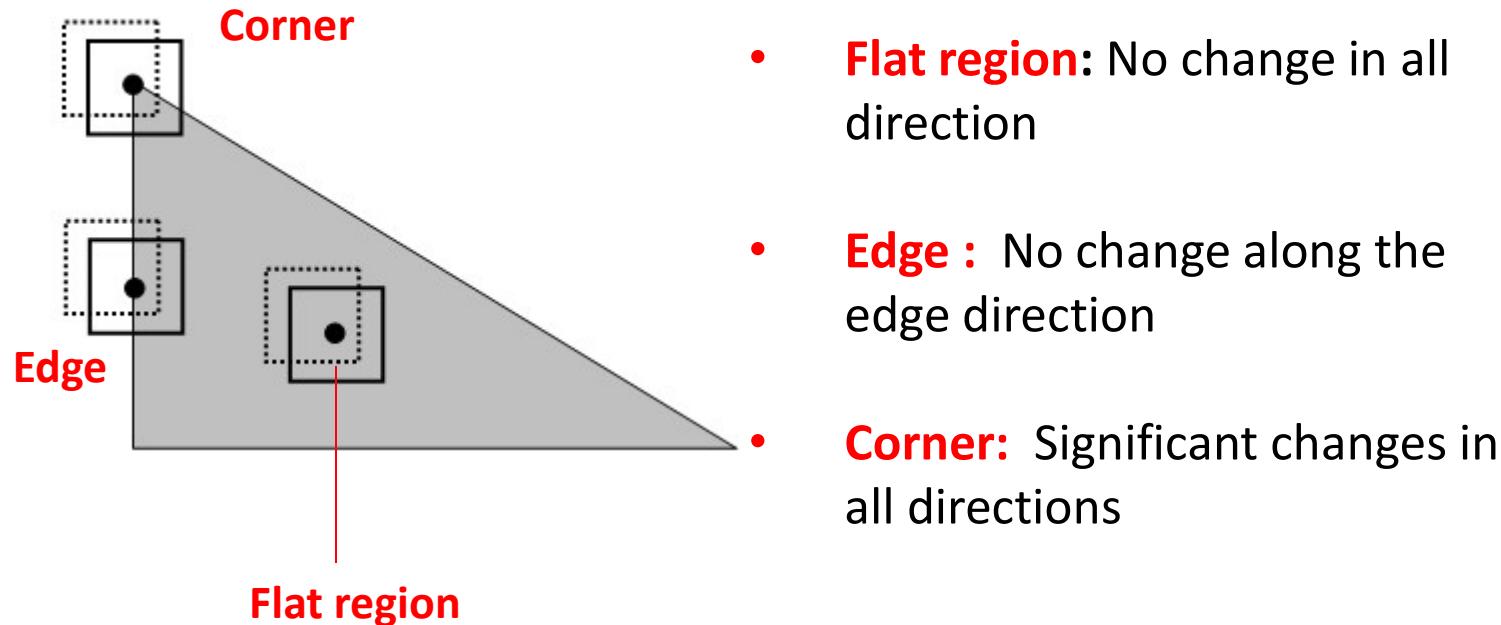
How do we define corner mathematically?



Harris, Chris, and Mike Stephens. "A combined corner and edge detector." Alvey vision conference. Vol. 15. 1988.

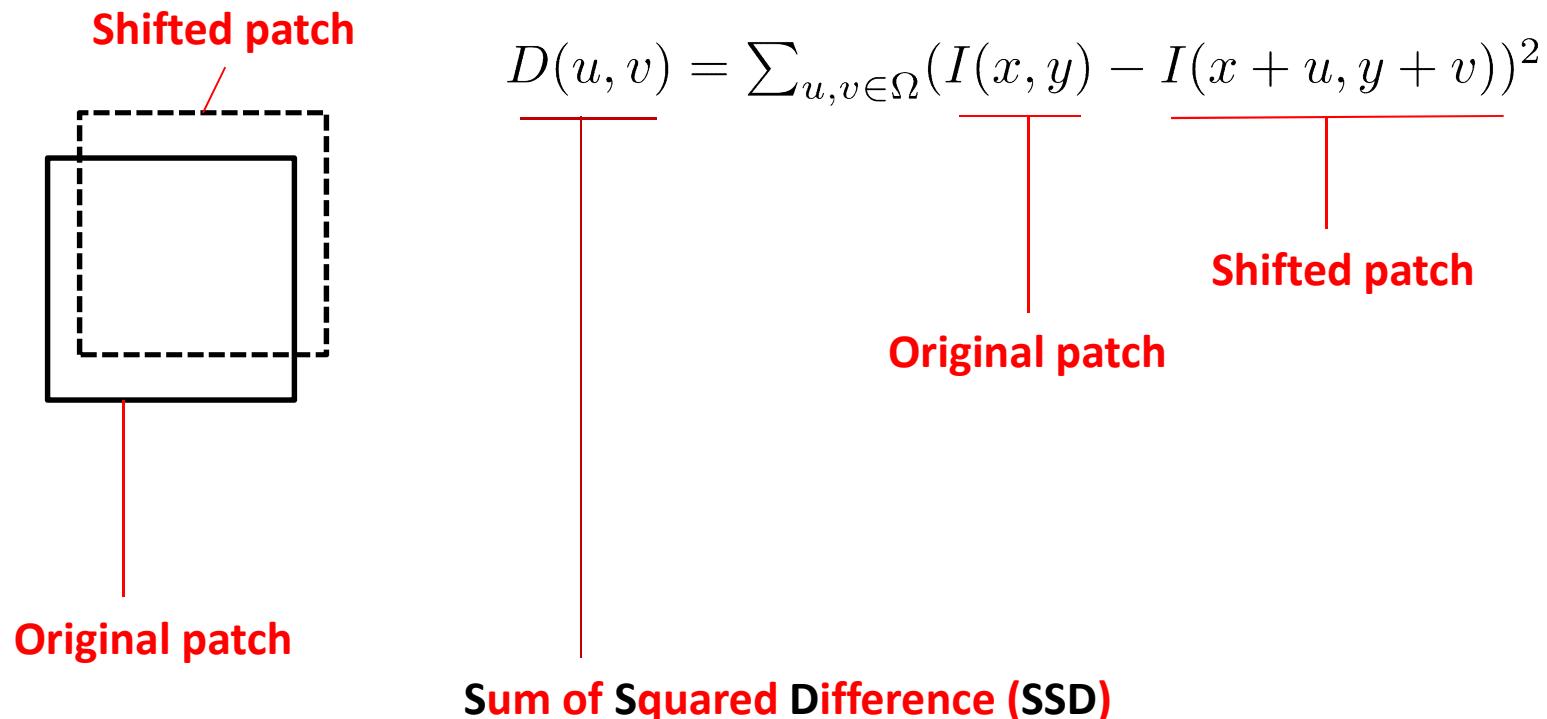
Feature detection

- Basic idea :
 - move the rectangle slightly around the original position and check the changes



Feature detection

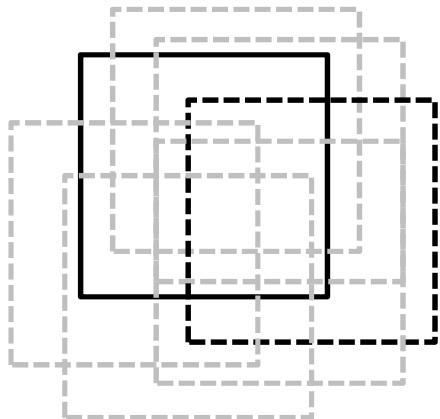
- Comparing the original image patch and the shifted image patch can be mathematically written as



Feature detection

- Maravec corner detector

Compare all nearby patches one by one to get the minimum SSD $D(u, v)$ as the corner response.



Problems:

1. high computational cost
2. Threshold of SSD is difficult to be set
(depends on the image content)

Moravec H P. Obstacle avoidance and navigation in the real world by a seeing robot rover[R]. STANFORD UNIV CA DEPT OF COMPUTER SCIENCE, **1980**.

Feature detection

- Harris corner detector:
- Using first-order Taylor expansion, we get

$$I(x + \delta x, y + \delta y) \approx I(x, y) + \nabla I(x, y) \begin{bmatrix} \delta x \\ \delta y \end{bmatrix} = I(x, y) + [I_x, I_y] \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

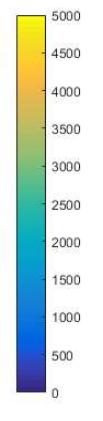
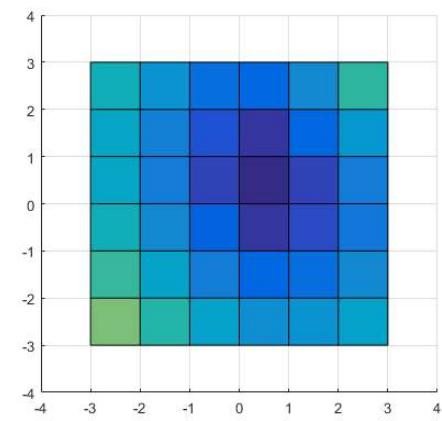
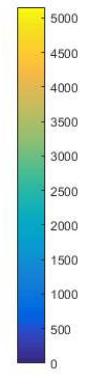
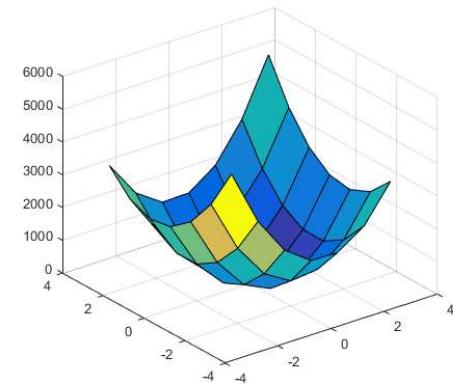
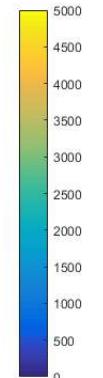
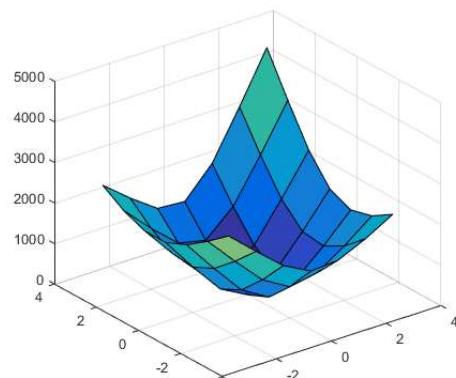
- Therefore, we can approximate the SSD function as

$$D(\delta x, \delta y) = \sum_{u,v \in \Omega} (I(x, y) - I(x + \delta x, y + \delta y))^2$$

$$\approx [\delta x, \delta y] \begin{pmatrix} \sum_{u,v \in \Omega} I_x(\delta x, \delta y) I_x(\delta x, \delta y) & \sum_{u,v \in \Omega} I_x(\delta x, \delta y) I_y(\delta x, \delta y) \\ \sum_{u,v \in \Omega} I_x(\delta x, \delta y) I_y(\delta x, \delta y) & \sum_{u,v \in \Omega} I_y(\delta x, \delta y) I_y(\delta x, \delta y) \end{pmatrix} \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

$$D(\delta x, \delta y) = [\delta x, \delta y] M \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

Feature detection



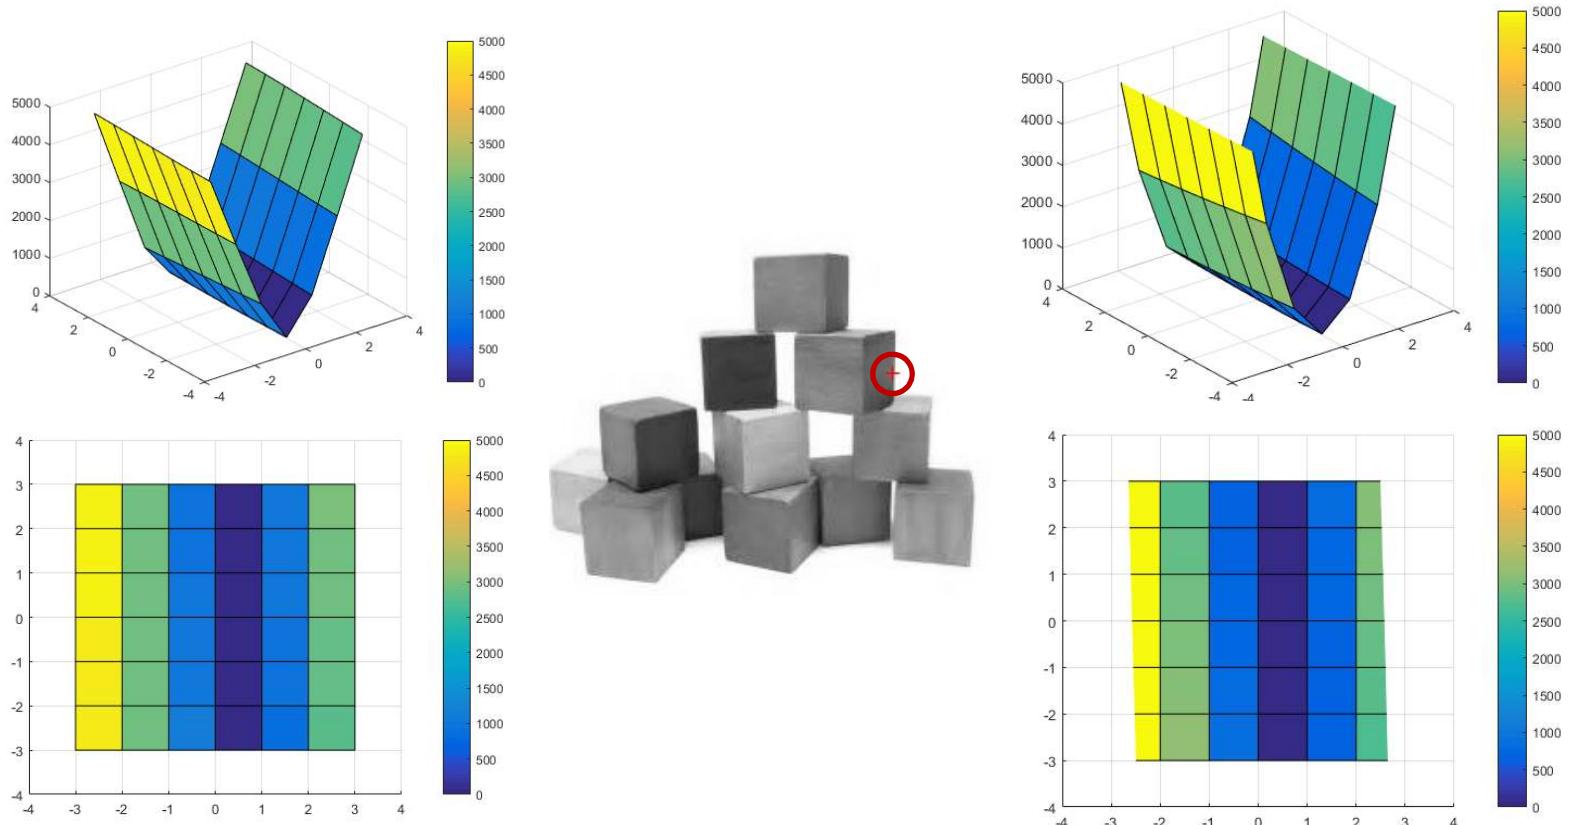
Original SSD function

$$D(\delta x, \delta y) = \sum_{\delta x, \delta y \in \Omega} (I(x, y) - I(x + \delta x, y + \delta y))^2$$

Approximated SSD function

$$\tilde{D}(\delta x, \delta y) = [\delta x, \delta y] M \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

Feature detection



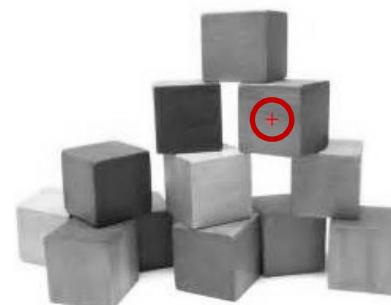
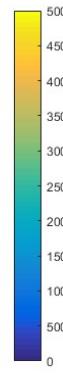
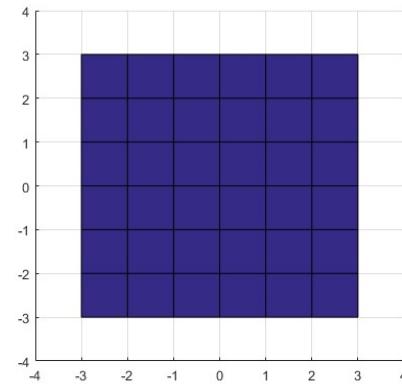
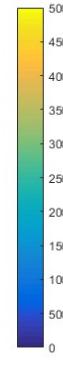
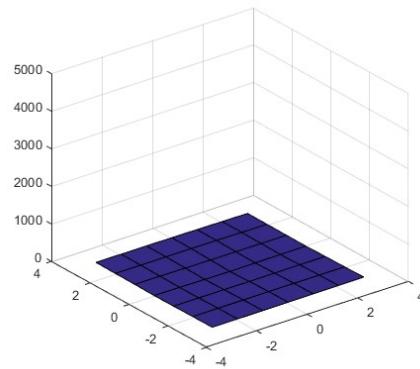
Original SSD function

$$D(\delta x, \delta y) = \sum_{\delta x, \delta y \in \Omega} (I(x, y) - I(x + \delta x, y + \delta y))^2$$

Approximated SSD function

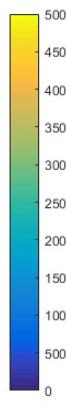
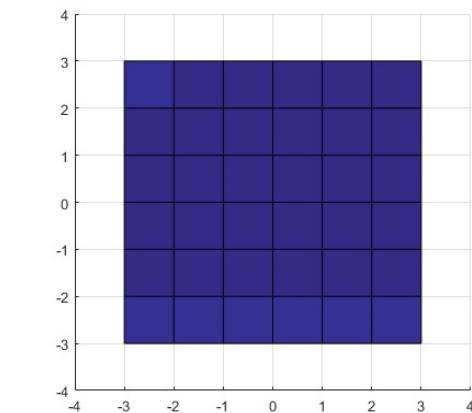
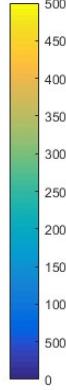
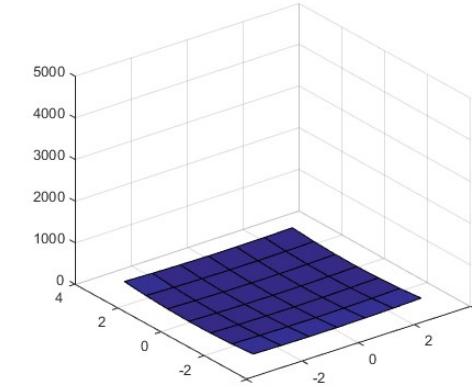
$$\tilde{D}(\delta x, \delta y) = [\delta x, \delta y] M \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

Feature detection



Original SSD function

$$D(\delta x, \delta y) = \sum_{\delta x, \delta y \in \Omega} (I(x, y) - I(x + \delta x, y + \delta y))^2$$



Approximated SSD function

$$\tilde{D}(\delta x, \delta y) = [\delta x, \delta y] M \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

Feature detection

- Approximated SSD function is very close to the real one locally.
- The auto-correlation matrix M plays a key role to detect the corners.

$$\tilde{D}(\delta x, \delta y) = [\delta x, \delta y] M \begin{bmatrix} \delta x \\ \delta y \end{bmatrix}$$

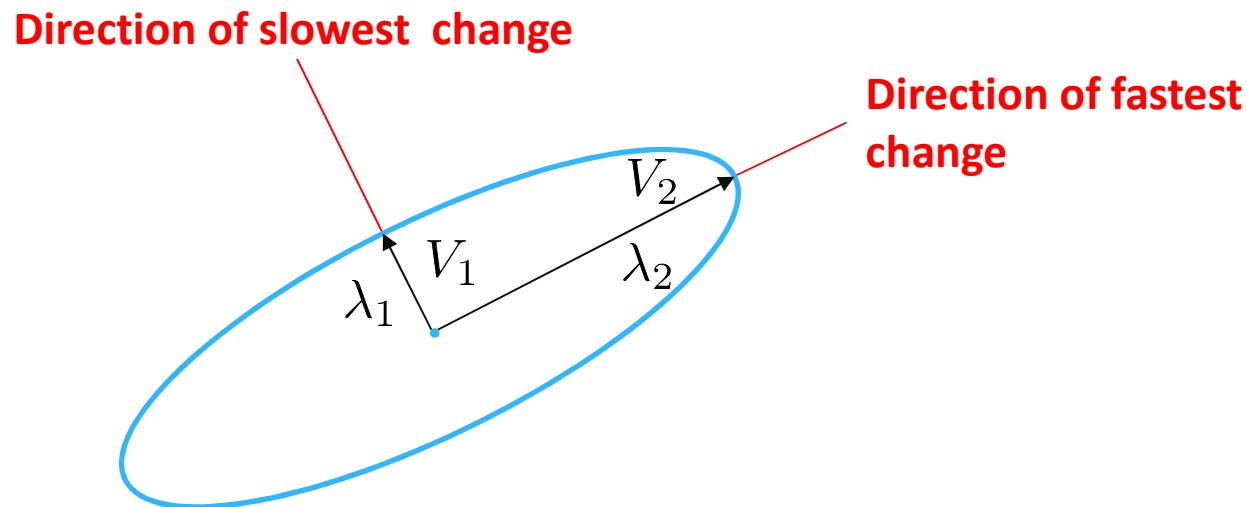
- M is a 2×2 positive definite symmetric matrix.
- Let $\lambda_1, \lambda_2, V_1, V_2$ be its eigenvalues and the corresponding eigenvectors.

$$MV_1 = \lambda_1 V_1$$

$$MV_2 = \lambda_2 V_2$$

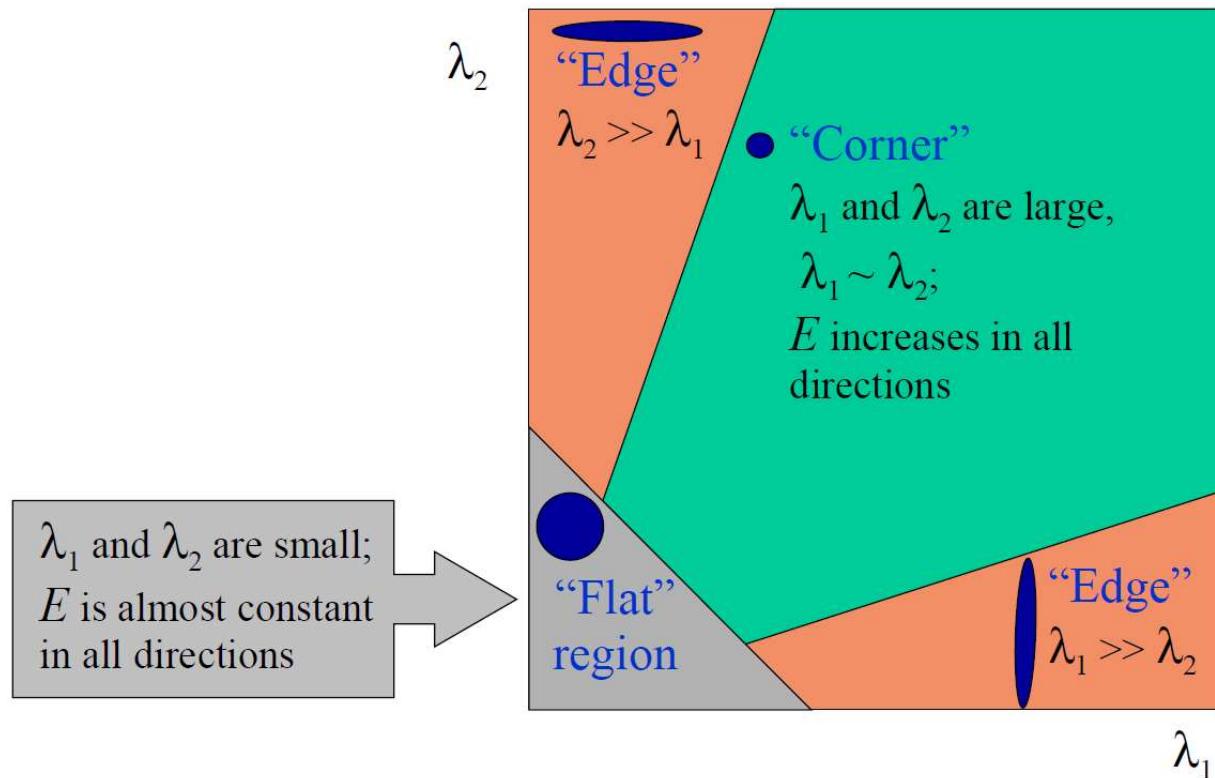
Feature detection

- Eigenvalue indicates the change of SSD



Feature detection

- Classify image points using eigenvalues of M:



Feature detection

- Measure of corner response
 - Harris method

$$r = \det(M) - k(\text{trace}(M))^2$$

$$\det(M) = \lambda_1$$

$$\text{trace}(M) = \lambda_1 + \lambda_2$$

k is an empirical constant (0.04 ~ 0.06)

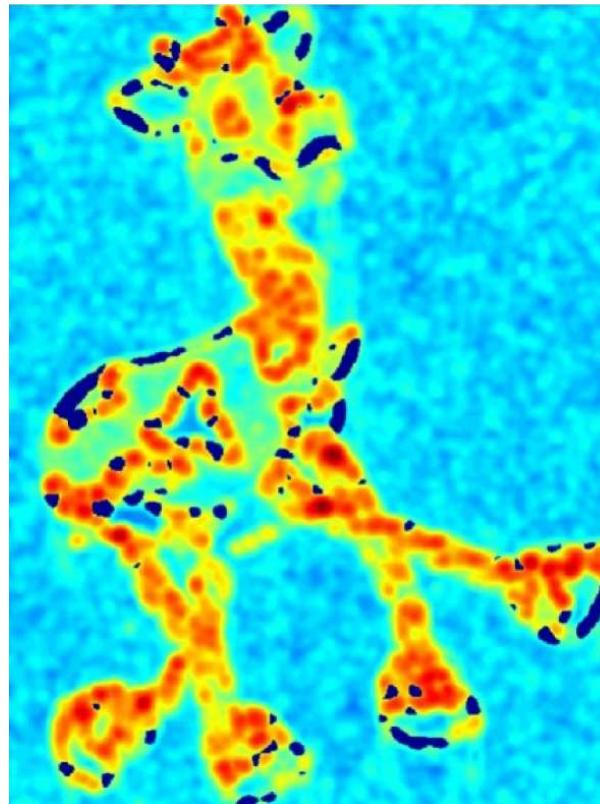
- Shi-Tomasi method

$$r = \min(\lambda_1, \lambda_2)$$

Shi, Jianbo, and Carlo Tomasi. "Good features to track." *Computer Vision and Pattern Recognition, 1994. Proceedings CVPR'94., 1994 IEEE Computer Society Conference on.* IEEE, 1994.

Feature detection

- Harris detector : work flow

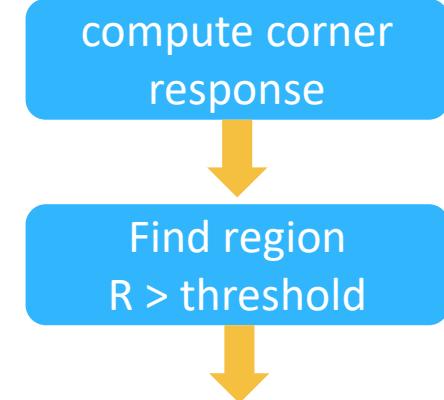
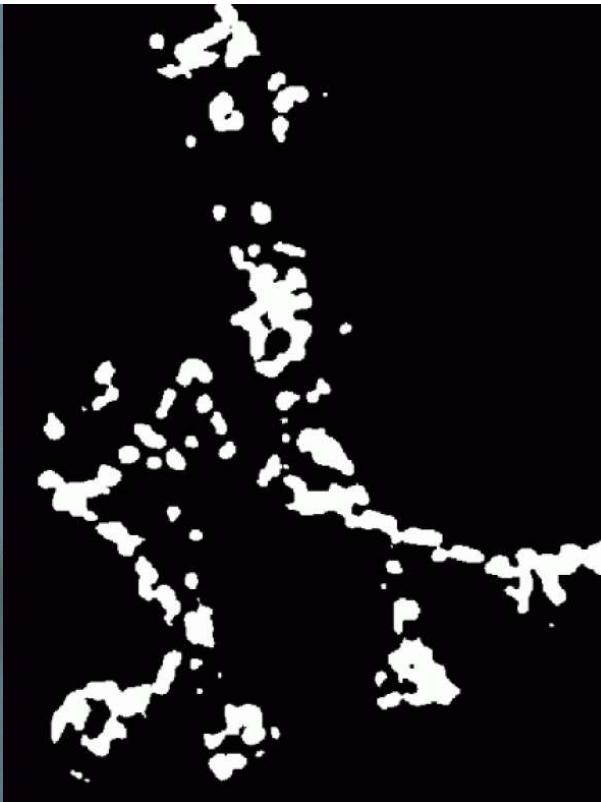


compute corner
response



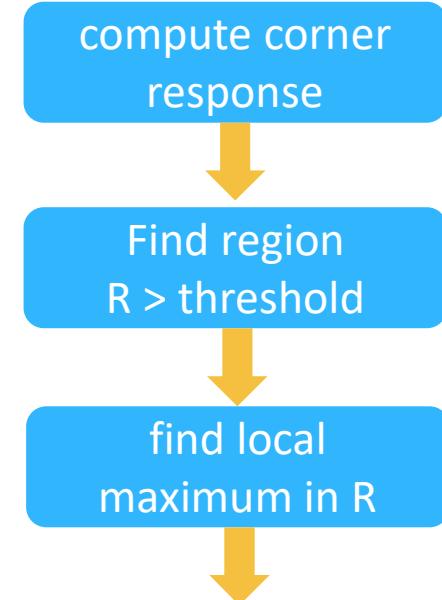
Feature detection

- Harris detector : work flow



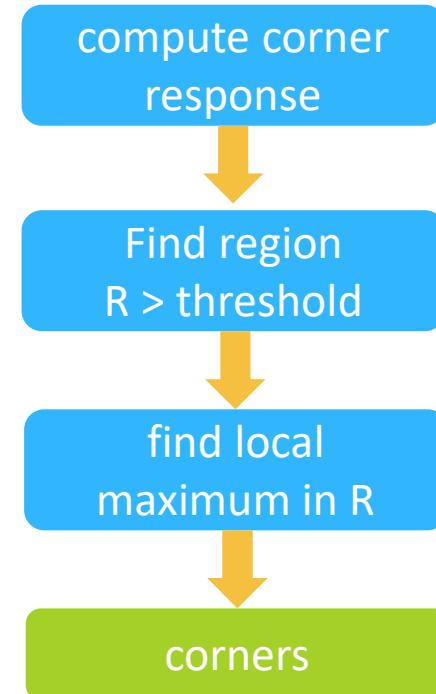
Feature detection

- Harris detector : work flow



Feature detection

- Harris detector : work flow



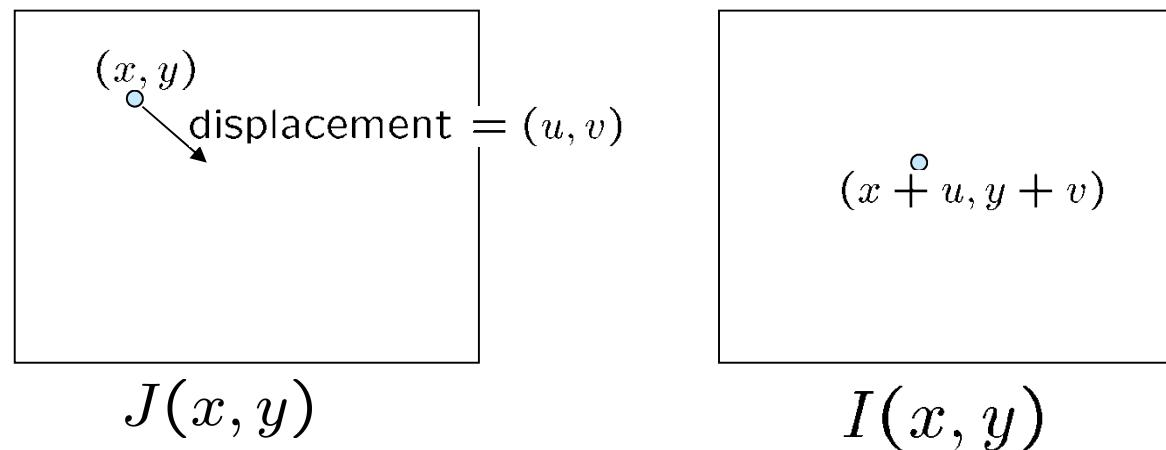
Summary

- Feature points should be recognizable, stable and accurate located.
- Harris corner is one kind of features that located on the position whose patch is very different from neighbor patches.
- The quality of corner is determined by the auto-correlation matrix M .
- The point with an auto-correlation matrix with two large Eigen values is detected as a corner.



Feature matching

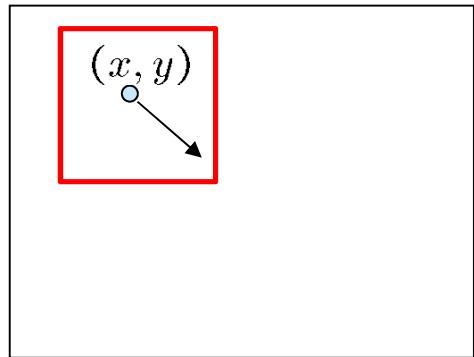
- Matching by tracking (for video sequence)
 - Lucas-Kanade optical flow algorithm



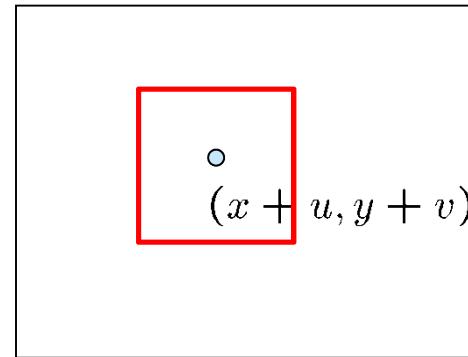
Given a feature point in the previous frame J , how to infer its position in the current frame I ?

Matching by tracking

- Assumption
 - Displacement is small.
 - Color has little change.



$J(x, y)$

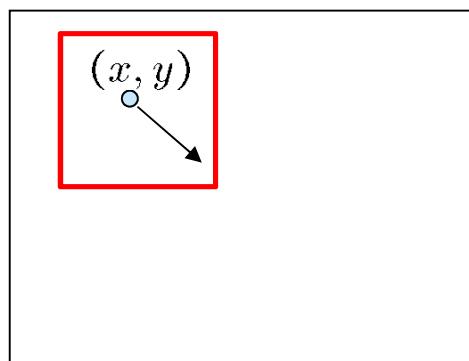


$I(x, y)$

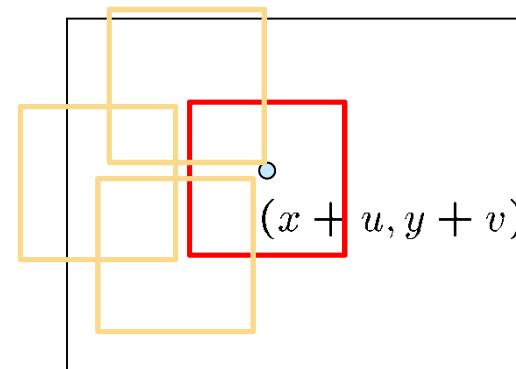
Compare color difference between two patches.

Matching by tracking

- A brute force searching method:
 - search the patch with the smallest color difference from original patch.



$J(x, y)$

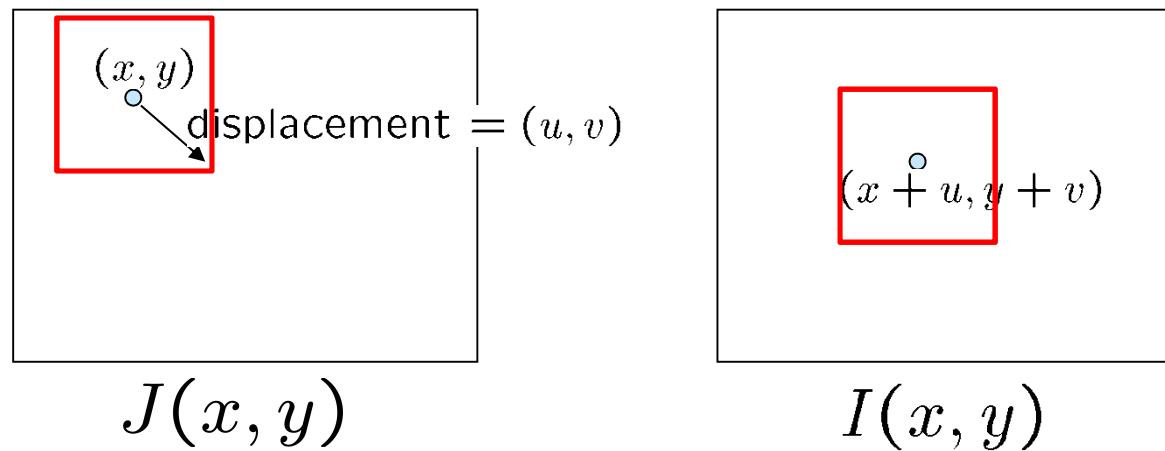


$I(x, y)$

- **Inefficient:** searching in a 10×10 window requires 100 times of patch comparison.
- **Inaccurate:** operates on pixels

Matching by tracking

- Lucas-Kanade optical flow method



- We are trying to minimize the **Sum of Squared Difference (SSD)**

$$D(u, v) = \sum_{u,v \in \Omega} (I(x + u, y + v) - J(x, y))^2$$

Matching by tracking

- Local minimum: let the gradient be zero!

$$\nabla D = \begin{bmatrix} \partial D / \partial u \\ \partial D / \partial v \end{bmatrix} = 0$$

$$D(u, v) = \sum_{u, v \in \Omega} (I(x + u, y + v) - J(x, y))^2$$

We can approximate this term by
Taylor series expansion (Lecture 2)

Matching by tracking

- First-order Taylor expansion

$$I(x+u, y+v) = I(x, y) + \nabla I(x, y) \begin{bmatrix} u \\ v \end{bmatrix} = I + I_x u + I_y v$$

$$S(u, v) = \sum_{x,y \in \Omega} (I(x+u, y+v) - J(x, y))^2$$

$$= \sum_{x,y \in \Omega} (I - J + I_x u + I_y v)^2$$

$$\begin{aligned} \nabla D &= \begin{bmatrix} \partial D / \partial u \\ \partial D / \partial v \end{bmatrix} \\ &= \begin{bmatrix} \sum_{x,y} I_x (I - J) + \sum_{x,y} I_x^2 u + \sum_{x,y} I_x I_y v \\ \sum_{x,y} I_y (I - J) + \sum_{x,y} I_x I_y u + \sum_{x,y} I_y^2 v \end{bmatrix} \end{aligned}$$

Matching by tracking

- Finally we get the following equation

$$\begin{bmatrix} \sum_{x,y} I_x^2 + \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y + \sum_{x,y} I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum_{x,y} I_x(I - J) \\ \sum_{x,y} I_y(I - J) \end{bmatrix}$$

Note that it is the auto-correlation matrix in Harris corner detection (P9).

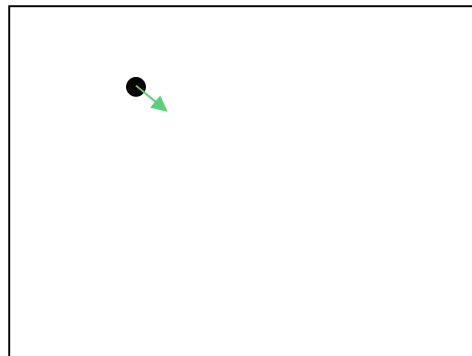
$$M\Delta X = b$$

$$\rightarrow \Delta X = M^{-1}\mathbf{b}$$

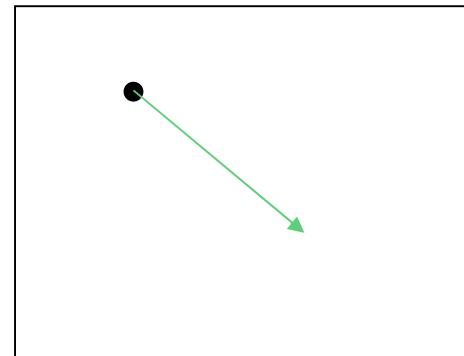
Finally, we get the displacement (flow)!

Matching by tracking

- Problem:
 - It assumes that the displacement is small (~ 1 pixel) that Taylor series expansion holds.
- How about **large** displacement value?



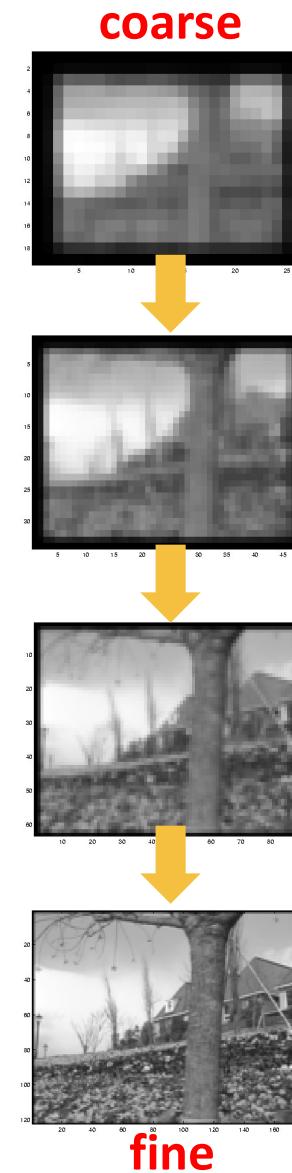
Small displacement



Large displacement

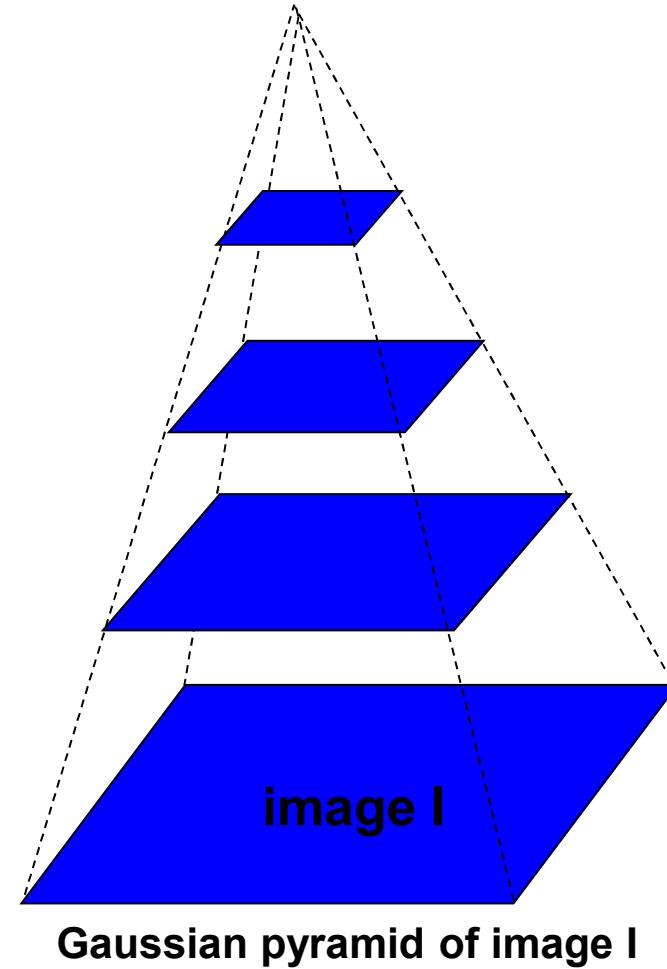
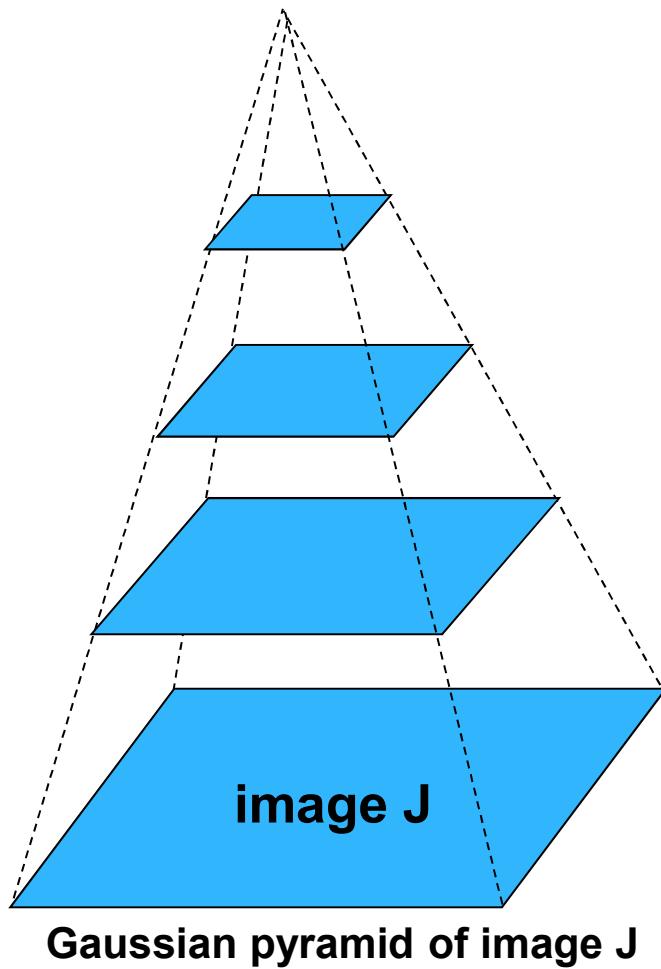
Matching by tracking

- Coarse-to-fine manner
 - Compute displacement in low resolution images (coarse)
 - Use the coarse displacement value as an initialization and refine the displacement in high resolution images(fine)
 - repeatedly run above steps until the finest level has been reached.



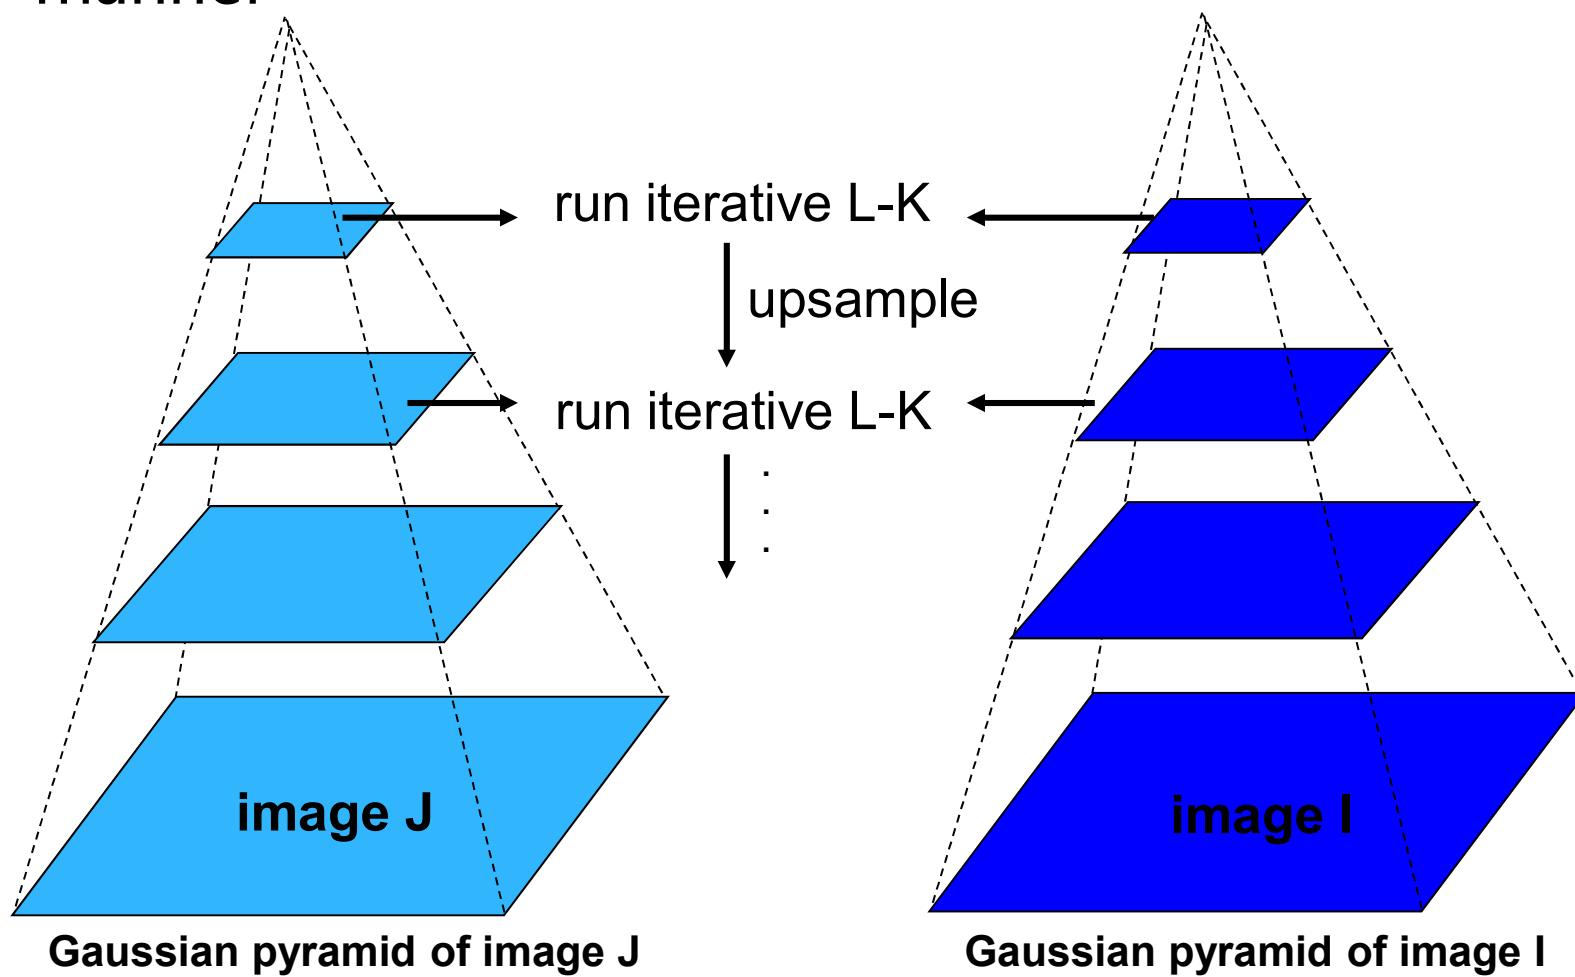
Matching by tracking

- Build image pyramid



Matching by tracking

- Run Lucas-Kanade algorithm in coarse-to-fine manner



Feature detection& matching

- OpenCV implementation

```
void goodFeaturesToTrack(InputArray image,  
                         OutputArray corners,  
                         int maxCorners,  
                         double qualityLevel,  
                         double minDistance,  
                         InputArray mask=noArray(),  
                         int blockSize=3,  
                         bool useHarrisDetector=false,  
                         double k=0.04 )
```

```
void calcOpticalFlowPyrLK(InputArray prevImg,  
                           InputArray nextImg,  
                           InputArray prevPts,  
                           InputOutputArray nextPts,  
                           OutputArray status,  
                           OutputArray err,  
                           Size winSize=Size(21,21),  
                           int maxLevel=3,  
                           TermCriteria criteria,  
                           int flags=0,  
                           double minEigThreshold=1e-4 )
```

Summary

- For video frames, we can use tracking to match the feature points in successive frames.
- When the displacement is small, matching becomes a problem of solving a linear equation.

$$M \Delta X = b$$

- If the displacement is large, we can use coarse-to-fine approach to compute displacement recursively.



Matching by feature descriptors

- Feature descriptor – a vector describes the appearance of the feature point.
- List of feature descriptors:

Spectra Descriptors: (Robust to viewpoint change or illumination change but slow!)

Scale Invariant Feature Transform (**SIFT**)

Speed-Up Robust Feature (**SURF**)

Histogram of Oriented Gradient (**HOG**)

Binary Descriptors: (suitable for real time processing)

Binary Robust Independent Elementary Features (**BRIEF**)

Oriented FAST and Rotated BRIEF (**ORB**)

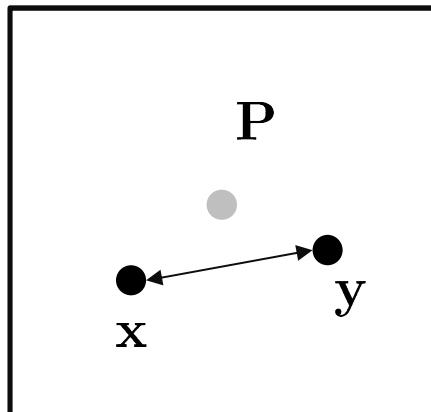
Binary Robust Invariance Scalable Key points (**BRISK**)

Fast Retina Key point (**FREAK**)

BRIEF descriptor

- Given a patch P , randomly sample two pixels x and y , compare their intensities by the following test:

$$\tau(P; x, y) = \begin{cases} 1 & \text{if } P(x) < P(y) \\ 0 & \text{otherwise} \end{cases}$$



Sample n_d pixel-pairs and run the binary test, we get the BRIEF descriptor of P as:

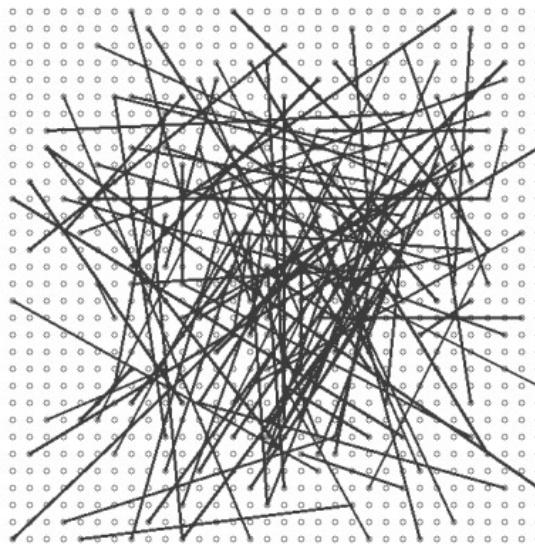
$$f_{n_d}(P) = \sum_{i=1}^{n_d} 2^{i-1} \tau(P; x, y)$$

$$n_d = 128, 256, 512$$

Note that the descriptor is a bit string.

BRIEF descriptor

- The approach to choosing the sample positions:



$$(\mathbf{x}, \mathbf{y}) \sim i.i.d. \quad \mathcal{N}(0, \frac{1}{25}S^2)$$

S is the width of the window.

BRIEF descriptor

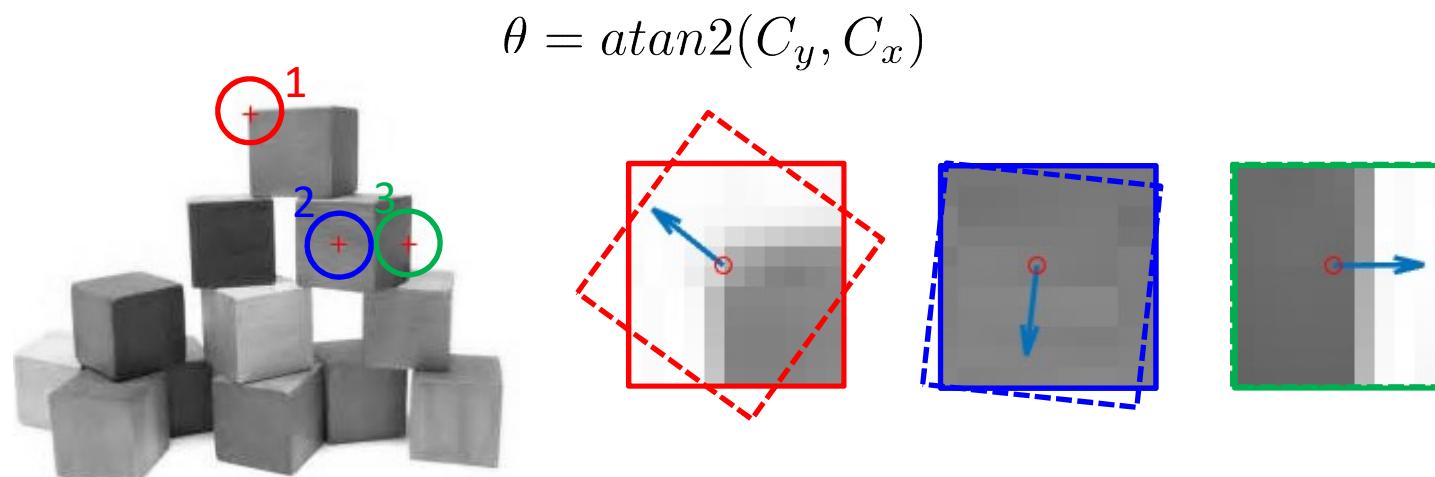
- **Hamming distance** to measure the distance between two descriptors.
- The Hamming distance is the number of positions at which the corresponding symbols are different:
 - "karolin" and "kathrin" is 3.
 - "karolin" and "kerstin" is 3.
 - **1011101** and **1001001** is 2.
 - **2173896** and **2233796** is 3.
- Very fast - only bit operations are involved.

ORB descriptor

- An improved version of BRIEF descriptor by adding **rotational invariance**.
- The corner orientation is defined by the intensity centroid:

$$C_x = \frac{\sum_{x,y} x I(x,y)}{\sum_{x,y} I(x,y)} \quad C_y = \frac{\sum_{x,y} y I(x,y)}{\sum_{x,y} I(x,y)}$$

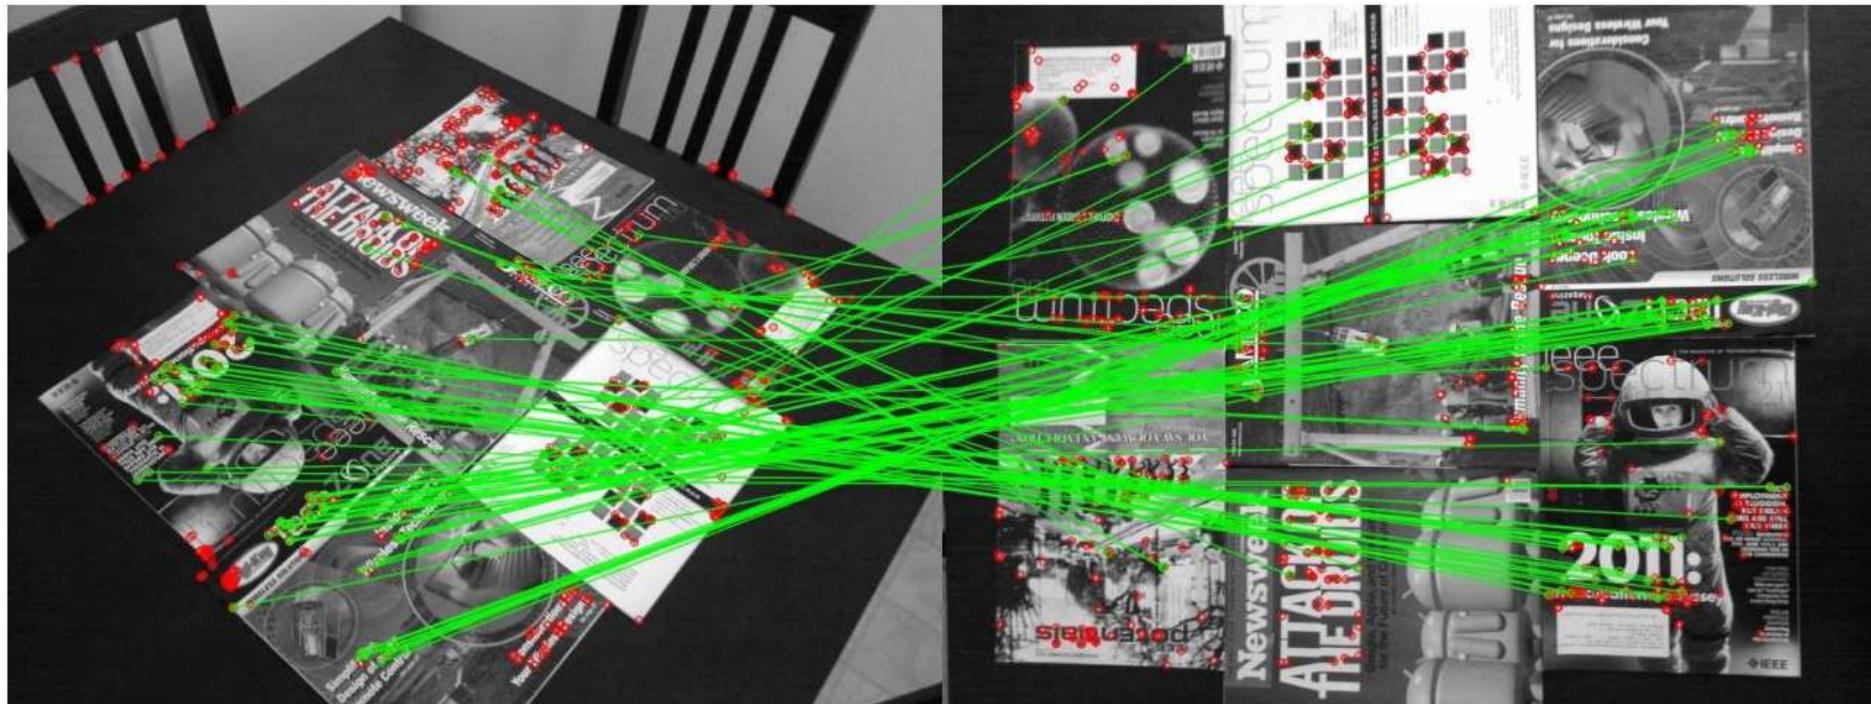
- The corner orientation is



ORB descriptor

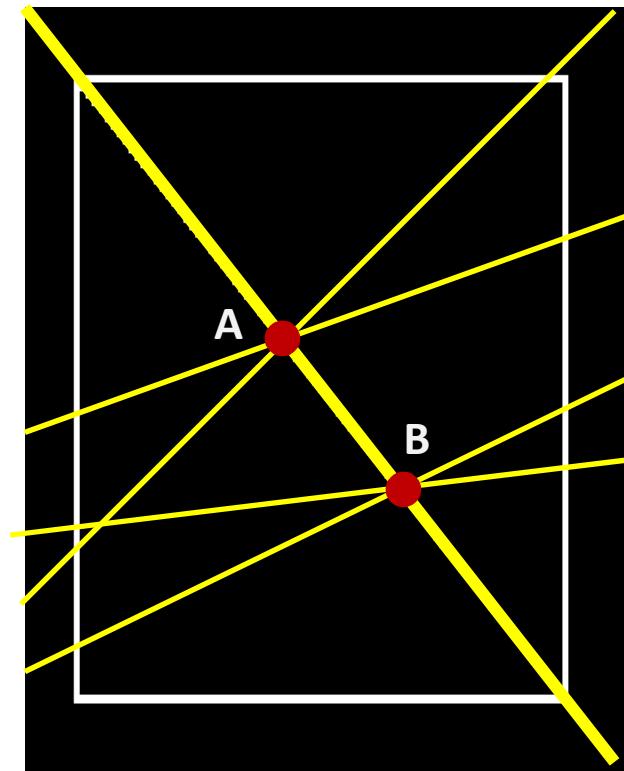
- ORB is much faster than other descriptors

Descriptor	ORB	SURF	SIFT
Time per frame (ms)	15.3	217.3	5228.7



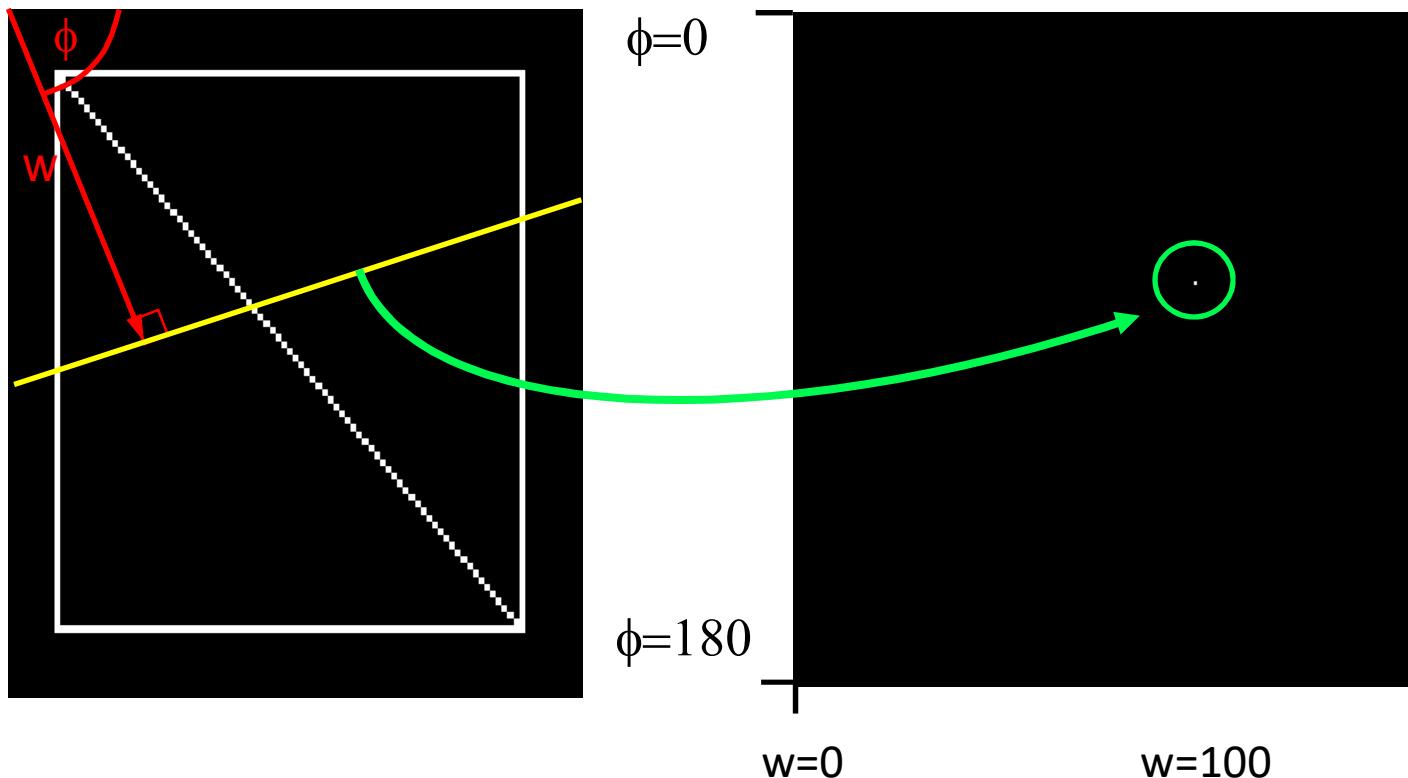
Line detection

- Hough transformation
 - The basic idea:
 - Each point votes for a group of lines that across this point.
 - The lines that has been voted by many points are straight lines.



Line detection

- A line can be represented by two parameters (ϕ, w) .
- Each white pixel corresponds to a point in the parameter space (Hough space) .



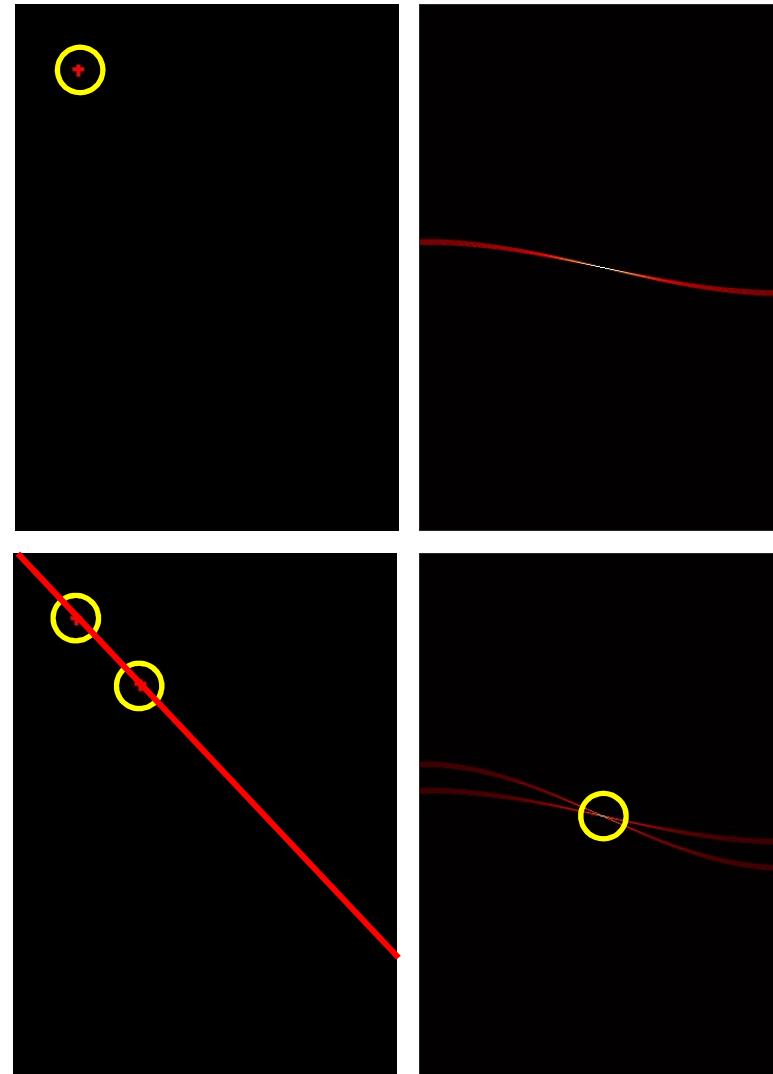
Line detection

One point in image space corresponds to a sinusoidal curve in Hough space

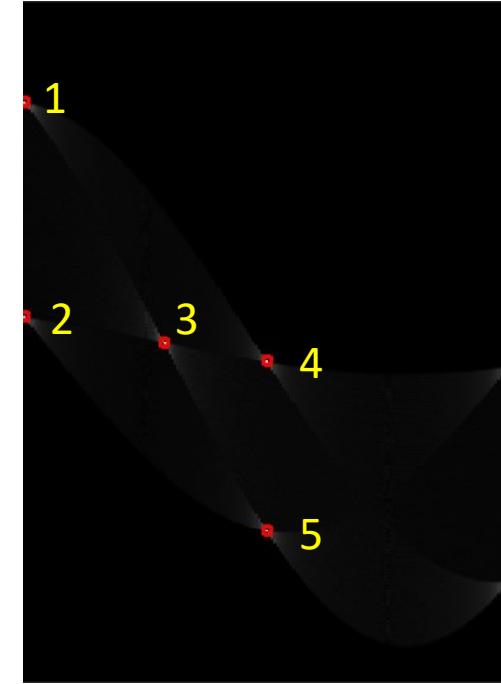
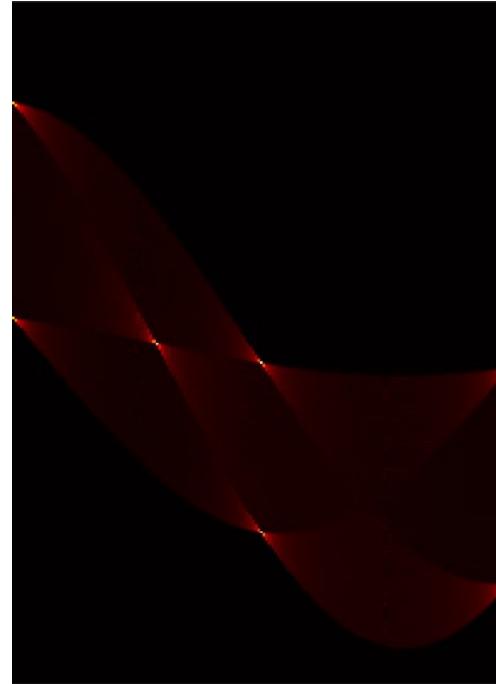
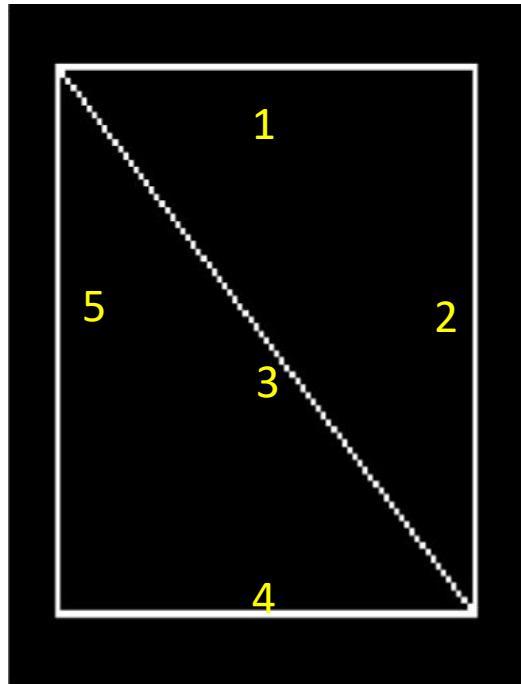
Two points correspond to two curves in Hough space

The intersection of those two curves has “two votes” .

This intersection represents the straight line in image space that passes through both points

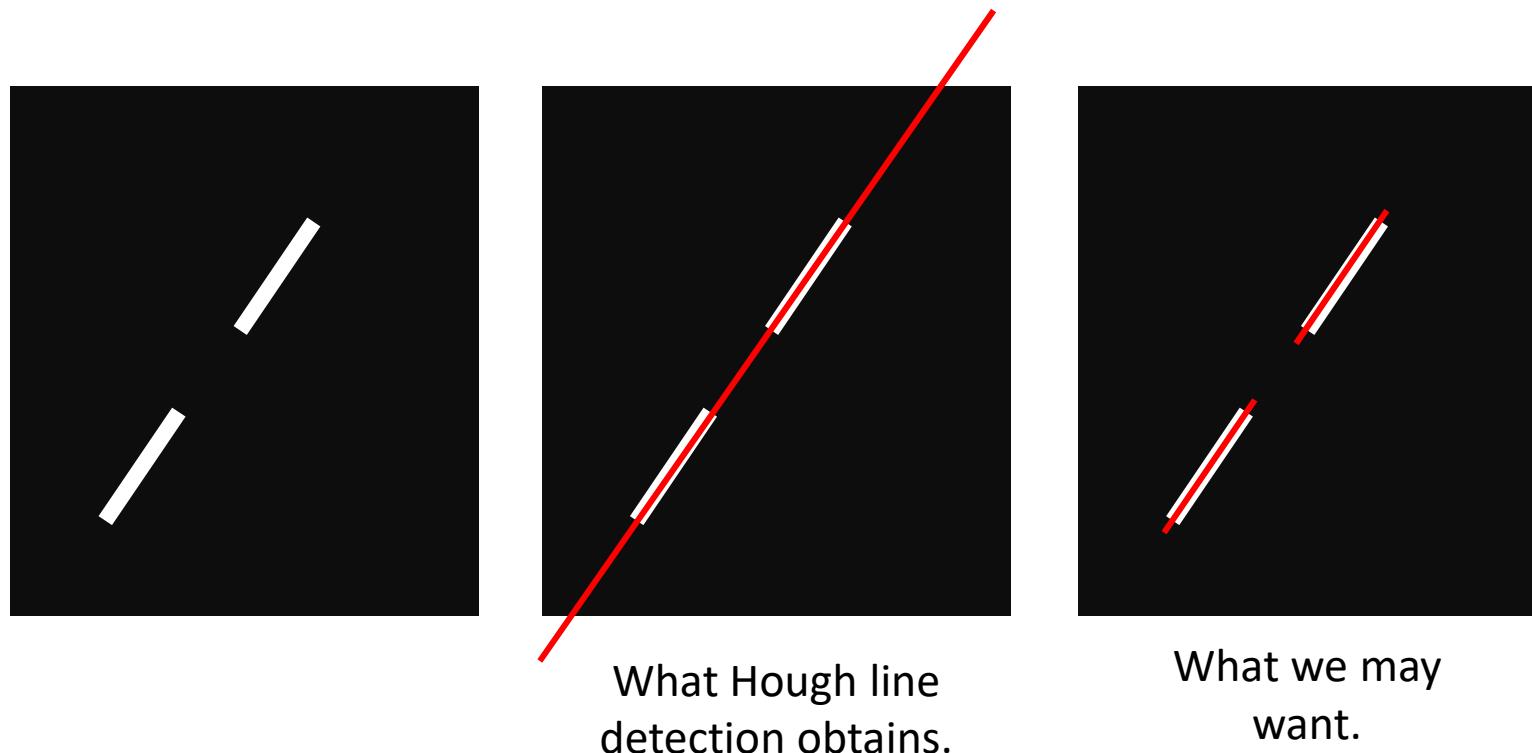


Line detection



Line detection

- One problem of Hough line detection is that it detects a line but not a line segment .
- In most situation, we want only line segments instead of infinite lines.



Line detection

- Other line detector
 - LSD line detector [1]



Von Gioi, R. G., Jakubowicz, J., Morel, J. M., & Randall, G. (2010). LSD: A fast line segment detector with a false detection control. *IEEE transactions on pattern analysis and machine intelligence*, 32(4), 722-732.

Summary

- BRIEF descriptor is a binary descriptor that compares a random pair of pixel intensities.
- ORB is an oriented version of BREIF descriptor by selecting the intensity centroid as the corner orientation.
- Hough transformation is transform a point into a sinuous curve in the parameter space (Hough space)
- Hough line detection detect only straight lines but not line segments.
- LSD line detector can be used to detect line segments