

Lecture 8.2

Structure from Motion

Thomas Opsahl

More-than-two-view geometry

Correspondences (matching)


- More views enables us to reveal and remove more mismatches than we can do in the two-view case
- More views also enables us to predict correspondences that can be tested with or without the use of descriptors

Scene geometry (structure)

- Effect of more views on determining the 3D structure of the scene?

Camera geometry (motion)


- Effect of more views on determining camera poses?



Structure from Motion

Problem

Given m images of n fixed 3D points, estimate the m projection matrices P_j and the n points X_j from the $m \cdot n$ correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{u}_{kj}$




Structure from Motion

Problem

Given m images of n fixed 3D points, estimate the m projection matrices P_j and the n points X_j from the $m \cdot n$ correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{u}_{kj}$


- We can solve for structure and motion when $2mn \geq 11m + 3n - 15$
- In the general/uncalibrated case, cameras and points can only be recovered up to a projective ambiguity ($\tilde{\mathbf{u}}_{ij} = P_i Q^{-1} Q \tilde{X}_j$)
- In the calibrated case, they can be recovered up to a similarity (scale)
 - Known as Euclidean/metric reconstruction




Structure from Motion

Problem


Given m images of n fixed 3D points, estimate the m projection matrices P_j and the n points X_j from the $m \cdot n$ correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{u}_{kj}$



- We can solve for structure and motion when
$$2mn \geq 11m + 3n - 15$$
- In the general/uncalibrated case, cameras and points can only be recovered up to a projective ambiguity ($\tilde{\mathbf{u}}_{ij} = P_i Q^{-1} Q \tilde{X}_j$)
- In the calibrated case, they can be recovered up to a similarity (scale)
 - Known as Euclidean/metric reconstruction



Projective reconstruction



Metric reconstruction

Images courtesy of Hartley & Zisserman <http://www.robots.ox.ac.uk/~vgg/hzbook/>

Structure from Motion

Problem


Given m images of n fixed 3D points, estimate the m projection matrices P_j and the n points X_j from the $m \cdot n$ correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{u}_{kj}$

- We can solve for structure and motion when
$$2mn \geq 11m + 3n - 15$$
- In the general/uncalibrated case, cameras and points can only be recovered up to a projective ambiguity ($\tilde{\mathbf{u}}_{ij} = P_i Q^{-1} Q \tilde{X}_j$)
- In the calibrated case, they can be recovered up to a similarity (scale)
 - Known as Euclidean/metric reconstruction

- This problem has been studied extensively and several different approaches have been suggested
- We will take a look at a couple of these
 - Sequential structure from motion
 - Bundle adjustment


Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow X_j$
 - Refine and extend the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation




Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation



Sequential SfM

- Initialize motion from two images
 - $F \rightarrow (P_1, P_2)$
 - $E \rightarrow (P_1, P_2) = (K_1[I \quad \mathbf{0}], K_2[{}^1R_2 \quad {}^1\mathbf{t}_2])$
- Initialize the 3D structure by triangulation
- For each additional view
 - Determine the projection matrix P_i , e.g. from 2D-3D correspondences $\mathbf{u}_{ij} \leftrightarrow \mathbf{X}_j$
 - Refine and extend the 3D structure by triangulation
- The resulting structure and motion can be refined in a process known as bundle adjustment




Bundle adjustment

- Non-linear method that refines structure and motion by minimizing the sum of squared reprojection errors


$$\epsilon = \sum_{i=1}^m \sum_{j=1}^n d(\tilde{\mathbf{u}}_{ij}, P_i \tilde{\mathbf{X}}_j)^2$$

- Camera calibration can be solved as part of bundle adjustment by including intrinsic parameters and skew parameters in the cost function
- Need initial estimates for all parameters!
 - 3 per 3D point
 - ~12 per camera depending on parameterization
 - Some intrinsic parameters, like the focal length, can be initialized from image EXIF data




Bundle adjustment

- There are several strategies that deals with the potentially extreme number of parameters
- Reduce the number of parameters by not including all the views and/or all the points
 - Perform bundle adjustment only on a subset and compute missing views/points based on the result
 - Divide views/points into several subsets which are bundle adjusted independently and merge the results
- Interleaved bundle adjustment
 - Alternate minimizing the reprojection error by varying only the cameras or only the points
 - This is viable since each point is estimated independently given fixed cameras, and similarly each camera is estimated independently from fixed points



Bundle adjustment

- Sparse bundle adjustment
 - For each iteration, iterative minimization methods need to determine a vector Δ of changes to be made in the parameter vector
 - In Levenberg-Marquardt each such step is determined from the equation
$$(J^T J + \lambda I) \Delta = -J^T \epsilon$$
where J is the Jacobian matrix of the cost function and ϵ is the vector of errors
 - For the bundle adjustment problem the Jacobian matrix has a sparse structure that can be exploited in computations



The sparse structure of the Jacobian matrix for a bundle adjustment problem with 3 cameras and 4 3D points

Figure courtesy of Hartley & Zisserman <http://www.robots.ox.ac.uk/~vgg/hzbook/>

Bundle adjustment

- Sparse bundle adjustment
 - For each iteration, iterative minimization methods need to determine a vector Δ of changes to be made in the parameter vector
 - In Levenberg-Marquardt each such step is determined from the equation
$$(J^T J + \lambda I) \Delta = -J^T \epsilon$$
where J is the Jacobian matrix of the cost function and ϵ is the vector of errors
 - For the bundle adjustment problem the Jacobian matrix has a sparse structure that can be exploited in computations
- Combined with parallel processing the before mentioned strategies has made it possible to solve extremely large SfM problems

Bundle adjustment

- Sparse bundle adjustment
 - For each iteration, iterative minimization methods need to determine a vector Δ of changes to be made in the parameter vector
 - In Levenberg-Marquardt each such step is determined from the equation
$$(J^T J + \lambda I) \Delta = -J^T \epsilon$$
where J is the Jacobian matrix of the cost function and ϵ is the vector of errors
 - For the bundle adjustment problem the Jacobian matrix has a sparse structure that can be exploited in computations
- Combined with parallel processing the before mentioned strategies has made it possible to solve extremely large SfM problems
- S. Agarwal et al, *Building Rome in a Day*, 2011
 - Cluster of 62-computers
 - 150 000 unorganized images from Rome
 - ~37 000 image registered
 - Total processing time ~21 hours
 - SfM time ~7 hours
- J. Heinly et al, *Reconstructing the World in Six Days*, 2015
 - 1 dual processor PC with 5 GPU's (CUDA)
 - ~96 000 000 unordered images spanning the globe
 - ~1.5 000 000 images registered
 - Total processing time ~5 days
 - SfM time ~17 hours

Bundle adjustment


- SBA – Sparse Bundle Adjustment
 - A generic sparse bundle adjustment C/C++ package based on the Levenberg-Marquardt algorithm
 - Code (C and Matlab mex) available at <http://www.ics.forth.gr/~lourakis/sba/>
 - CVSBA is an OpenCV wrapper for SBA www.uco.es/investiga/grupos/ava/node/39/
- Ceres
 - By Google (used in production since 2010)
 - A C++ library for modeling and solving large, complicated optimization problems like SfM
 - Homepage: www.ceres-solver.org
 - Code available on GitHub <https://github.com/ceres-solver/ceres-solver>
- GTSAM – Georgia Tech Smoothing and Mapping
 - A C++ library based on factor graphs that is well suited for SfM ++
 - Code (C++ library and Matlab toolbox) available at <https://borg.cc.gatech.edu/borg/download>
- g²o – General Graph Optimization
 - Open source C++ framework for optimizing graph-based nonlinear error functions
 - Homepage: <https://openslam.org/g2o.html>
 - Code available on GitHub <https://github.com/RainerKuemmerle/g2o>

Bundle adjustment

- Bundler
 - A structure from motion system for unordered image collections written in C and C++
 - SfM based on a modified version SBA (default) or Ceres
 - Homepage: <http://www.cs.cornell.edu/~snavely/bundler/>
 - Code available on GitHub
https://github.com/snavely/bundler_sf
- RealityCapture
 - A state-of-the-art photogrammetry software that automatically extracts accurate 3D models from images, laser-scans and other input
 - Homepage: <https://www.capturingreality.com/>
- VisualSfM
 - A GUI application for 3D reconstruction using structure from motion
 - Output works with other tools that performs dense 3D reconstruction
 - Homepage: <http://ccwu.me/vsfm/>


Example

Holmenkollen 2-view SfM




Example

Holmenkollen 2-view SfM




Example

Holmenkollen 2-view SfM




Example

Holmenkollen 3-view SfM




Example

Holmenkollen 4-view SfM




Example

Holmenkollen 110-view SfM




Example

Holmenkollen 110-view SfM



Summary

- Structure from motion
 - Sequential SfM
 - Bundle adjustment
- Additional reading:
 - Szeliski: 7.3-7.5
- Optional reading:
 - Snavely N. Seitz S. M., Szeliski R., *Modeling the World from Internet Photo Collections*, 2007
 - S. Agarwal et al, *Building Rome in a Day*, 2011
 - J. Heinly et al, *Reconstructing the World in Six Days*, 2015



$$\epsilon = \sum_{i=1}^m \sum_{j=1}^n d(\tilde{u}_{ij}, P_i \tilde{X}_j)^2$$

